

Intermediate Python & a Glimpse into AI applications

Class 5

Pepe Bonet Giner

18th January 2023

Index Class 5

Topic 1: Short recap

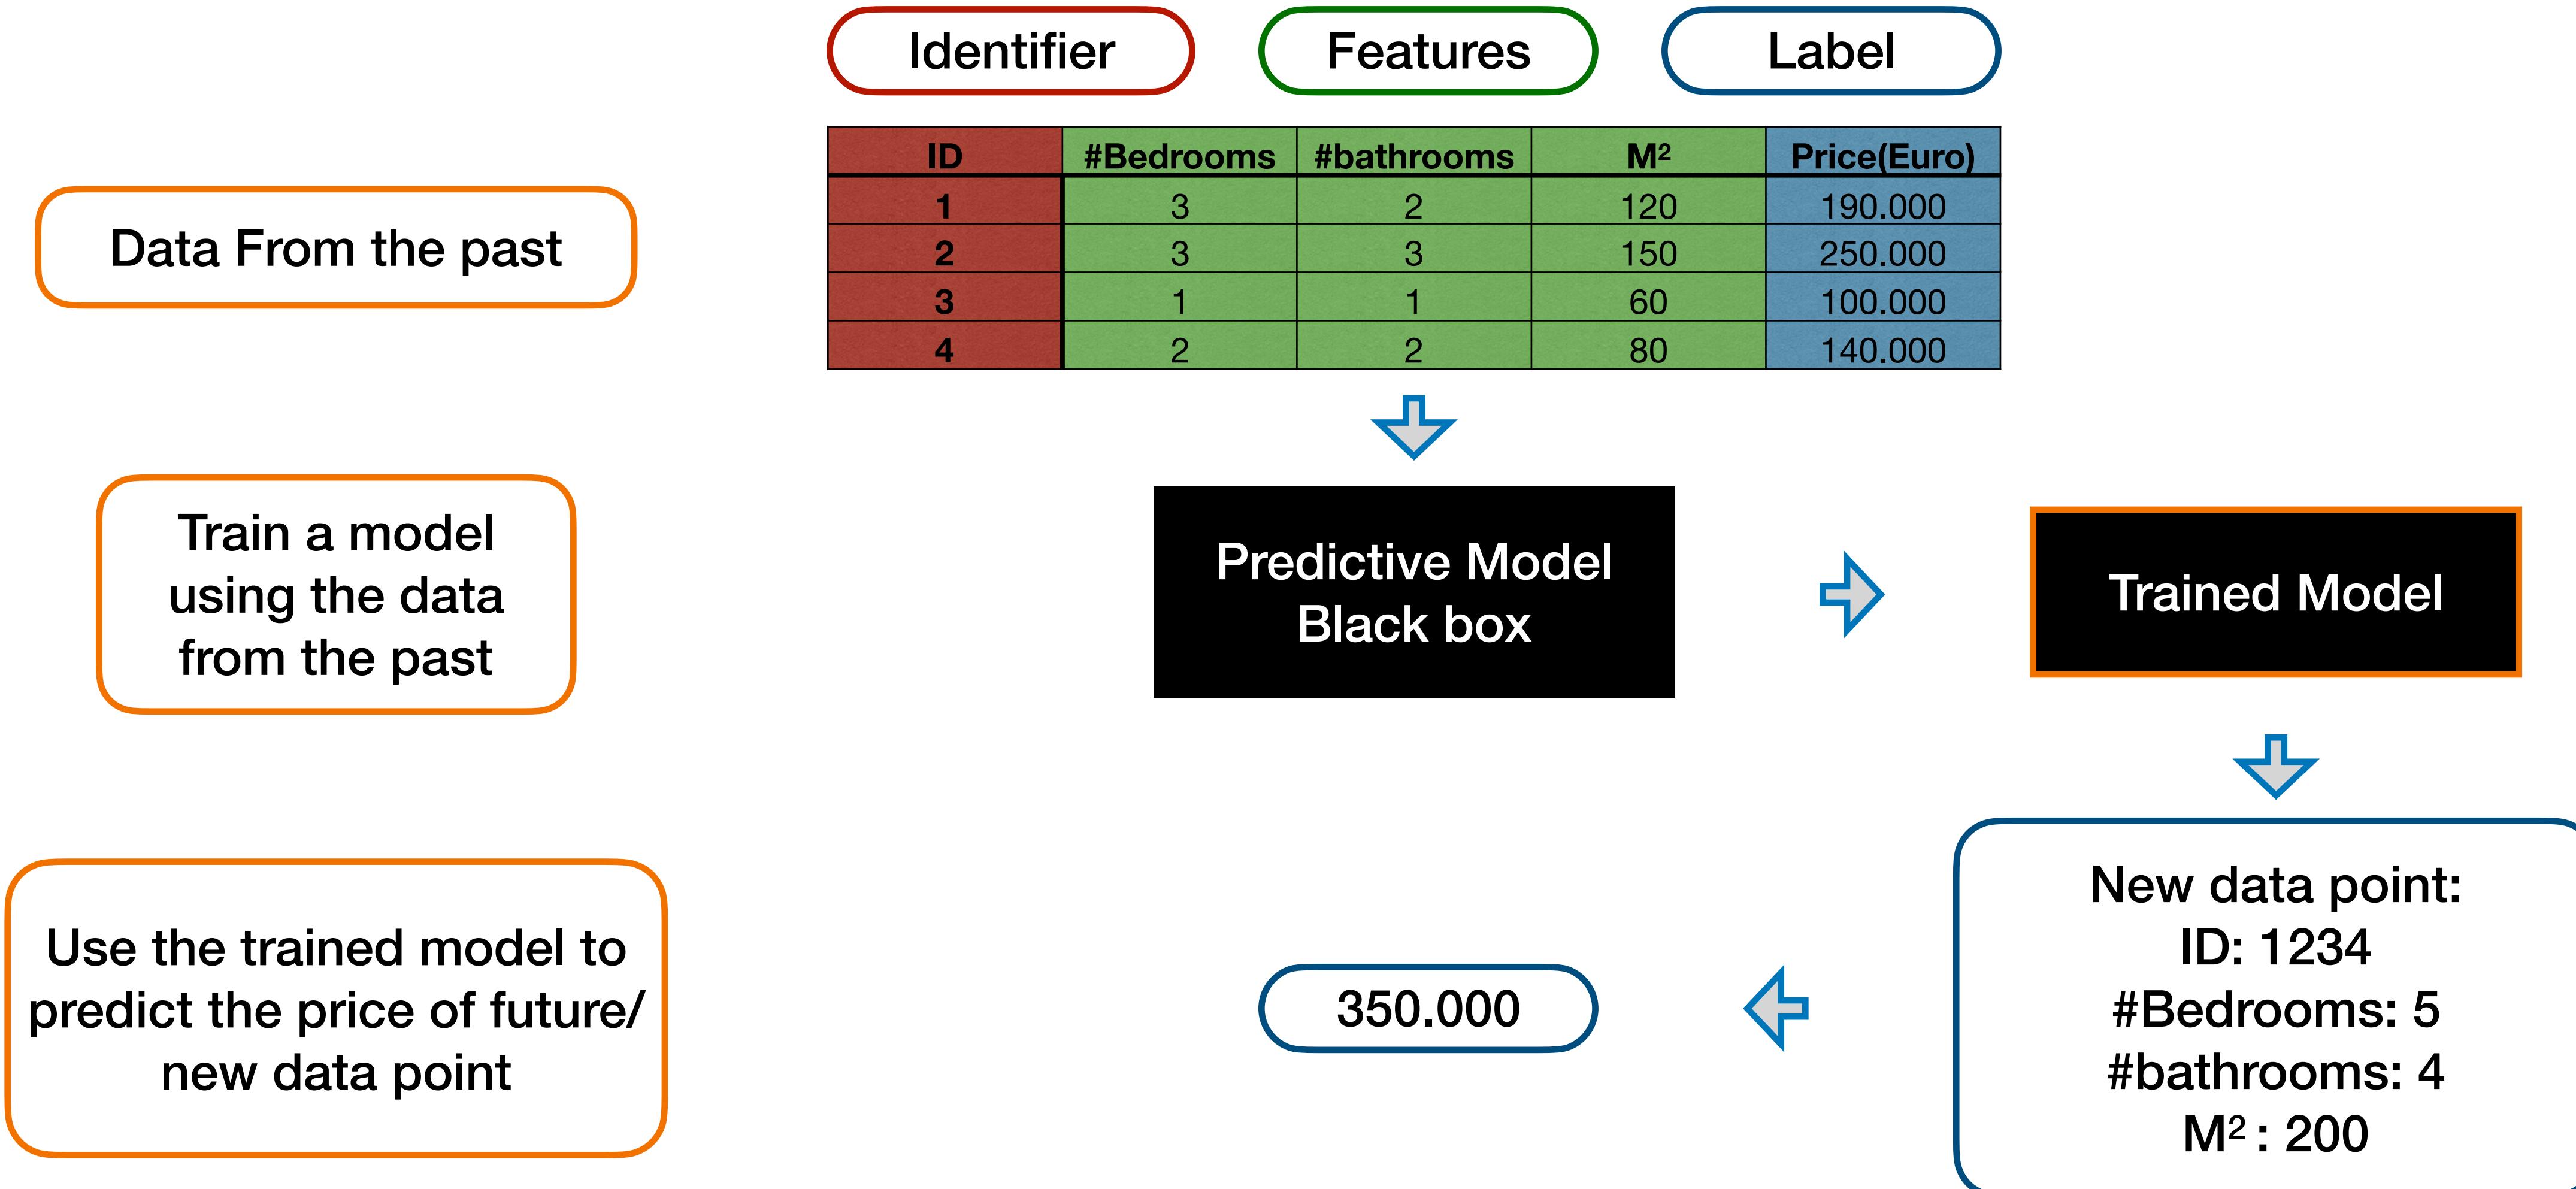
Topic 2: Intro to ML models in Python. Linear Regression

Topic 3: Scikit Learn models Overview

Topic 4: Final Project Example

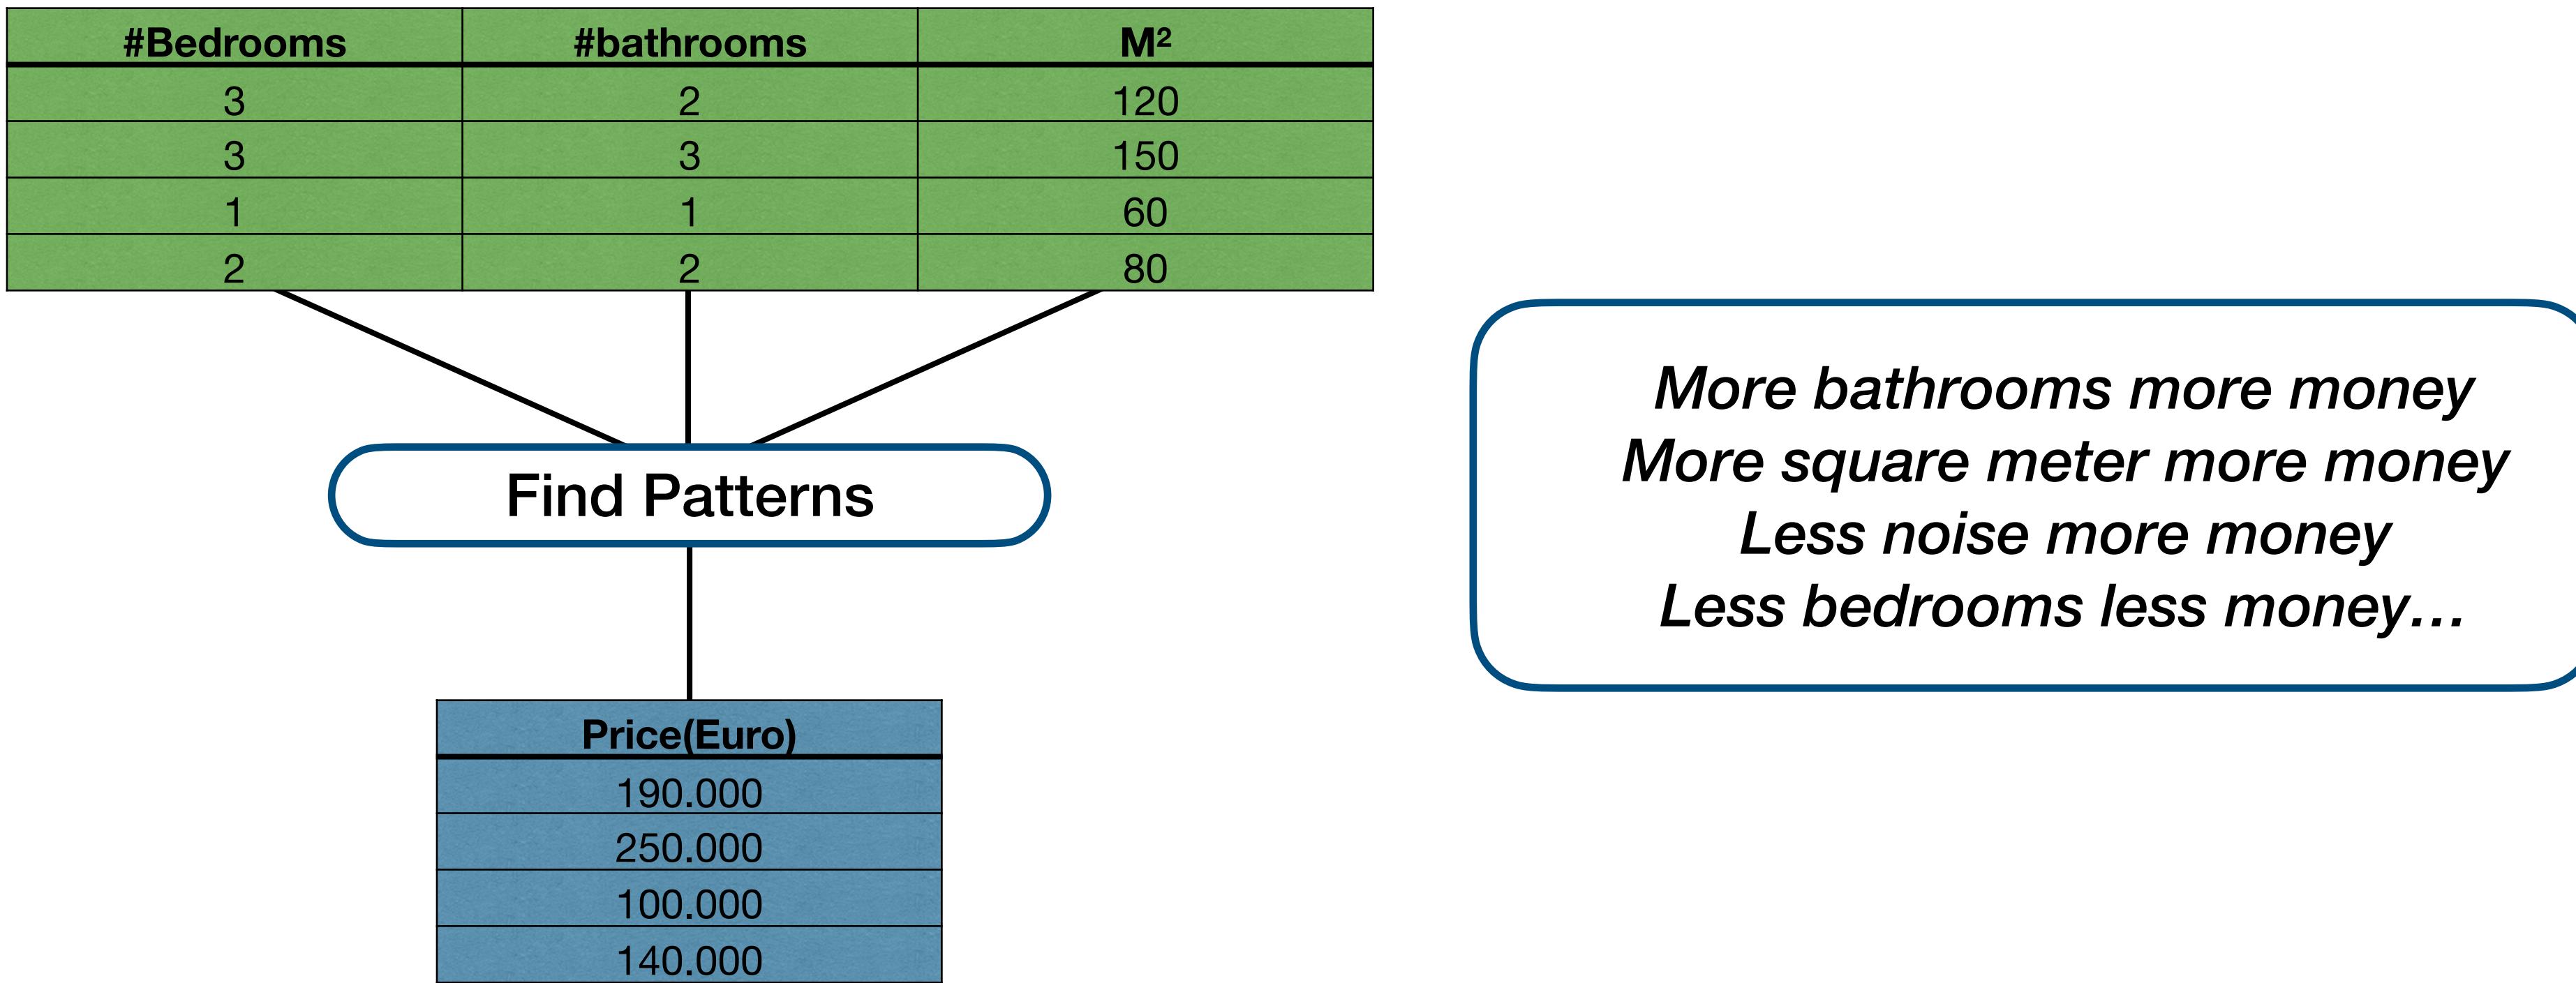
Topic 1: Recap

Recap: How do we learn from the past? We need something. A model



Recap: What is a model?

- Models are programs that can find patterns in the data that is fed to them



AI/ML. Classification & regression

ML: Data (input) + Answers (predictive label) → Rules (Model)

Input Data

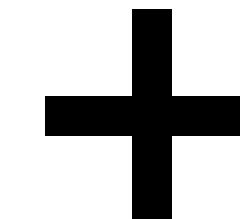
#Bedrooms	#bathrooms	M ²
3	2	120
3	3	150
1	1	60
2	2	80

Answers

Price(Euro)
190.000
250.000
100.000
140.000

Rules

Model



Fault

No
Yes
No
No

Classification

Price(Euro)

190.000
250.000
100.000
140.000

Regression

*Classification is about predicting a label
(Categorical) and regression is about predicting a quantity (Numerical)*

Computers only understand numbers!

Categorical Features

Gender	Parent Education	Age	Grade
Male	College	21	8
Female	High School	22	9
Male	None	20	6
Female	College	21	5

Numerical Features

Label Encoding

Parent Education
2
1
0
2



One Hot Encoding

Parent Education
College
High School
None
College

Computers need to code the categorical features/columns as numbers. How do we do it?

We have two main approaches:

- Label Encoding
- One Hot Encoding

Parent Ed_College	Parent Ed_High School	Parent Ed_None
1	0	0
0	1	0
0	0	1
1	0	0

Topic 2: Intro to ML models in Python.

Linear Regression

Train & Test sets. Why are they necessary?

Overfitting: The model is not good on never seen data

Underfitting: The model is bad everywhere

The reason to have a separate train and test set is to avoid overfitting of the model to the data

Train & Test sets. Why are they necessary?

Overfitting: The model is not good on never seen data

Underfitting: The model is bad everywhere

The reason to have a separate train and test set is to avoid overfitting of the model to the data

Training Set

Gender	Parent Education	Age	Grade
Male	College	21	8
Female	High School	22	9

Test Set

Gender	Parent Education	Age	Grade
Male	None	20	6
Female	College	21	5

To train the model. Can represent up to 90 % of the data

To test how the model generalizes to never seen data

Exercise

- Generate a train and test set of your dataset

Our Pandas for ML models is Scikit Learn



Scikit learn is an ML library for Python that features various classification, regression and clustering algorithms

It will allow us to build a model and make predictions in 4 lines of code

```
from sklearn import datasets, linear_model
from sklearn.metrics import mean_absolute_error

# Create linear regression object
regr = linear_model.LinearRegression()

# Train the model using the training sets
regr.fit(trainX, trainY)

# Make predictions using the testing set
y_pred = regr.predict(testX)

# The mean absolute error
print("MAE: %.2f" % mean_absolute_error(testY, y_pred))

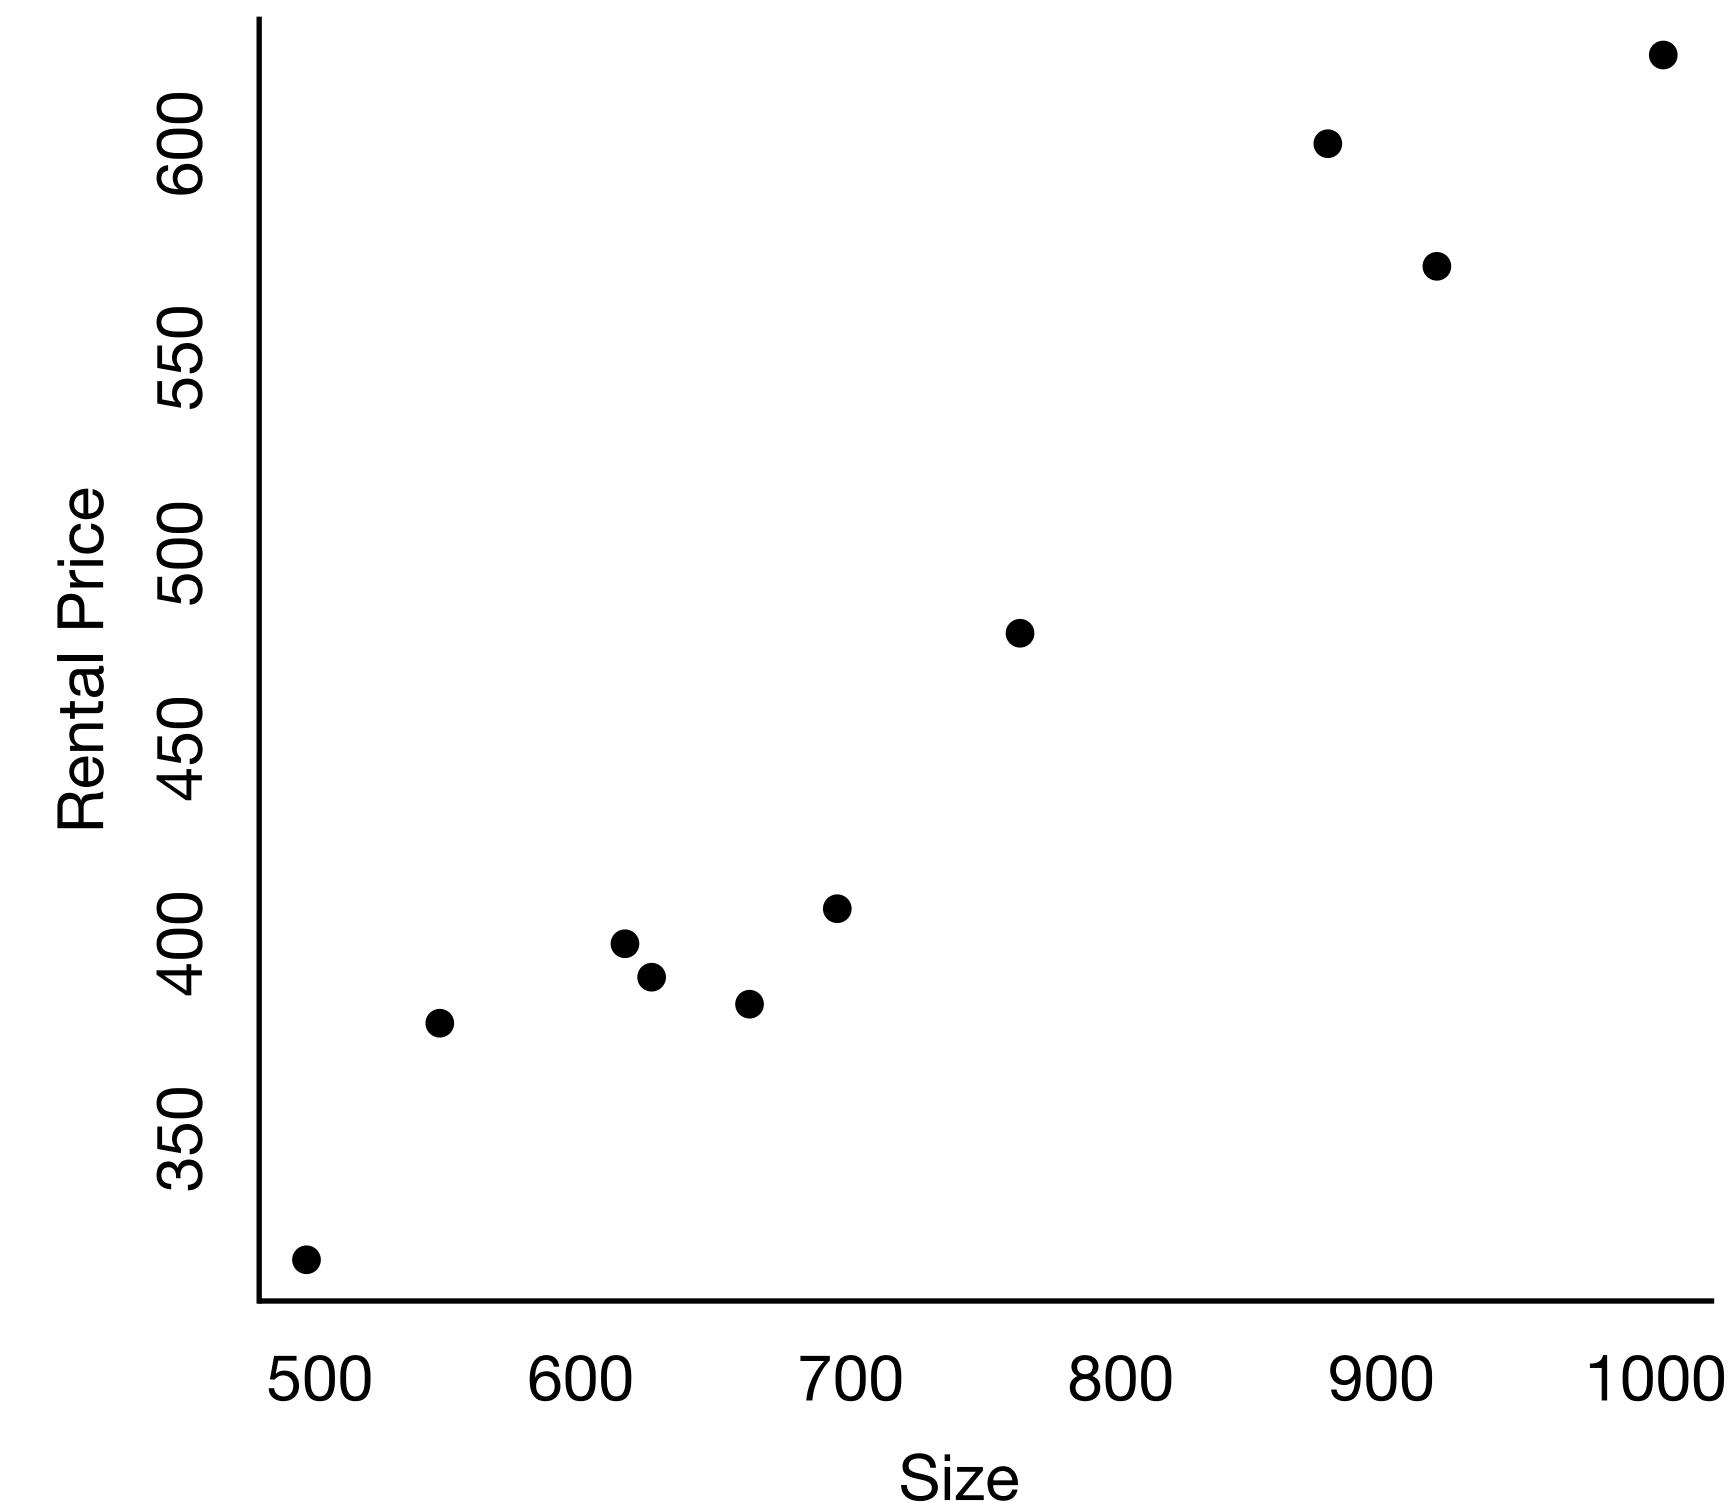
MAE: 4.41
```

Linear regression. Explained

ID	SIZE	FLOOR	BROADBAND RATE	ENERGY RATING	RENTAL PRICE
1	500	4	8	C	320
2	550	7	50	A	380
3	620	9	7	A	400
4	630	5	24	B	390
5	665	8	100	C	385
6	700	4	8	B	410
7	770	10	7	B	480
8	880	12	50	A	600
9	920	14	8	C	570
10	1,000	9	24	B	620

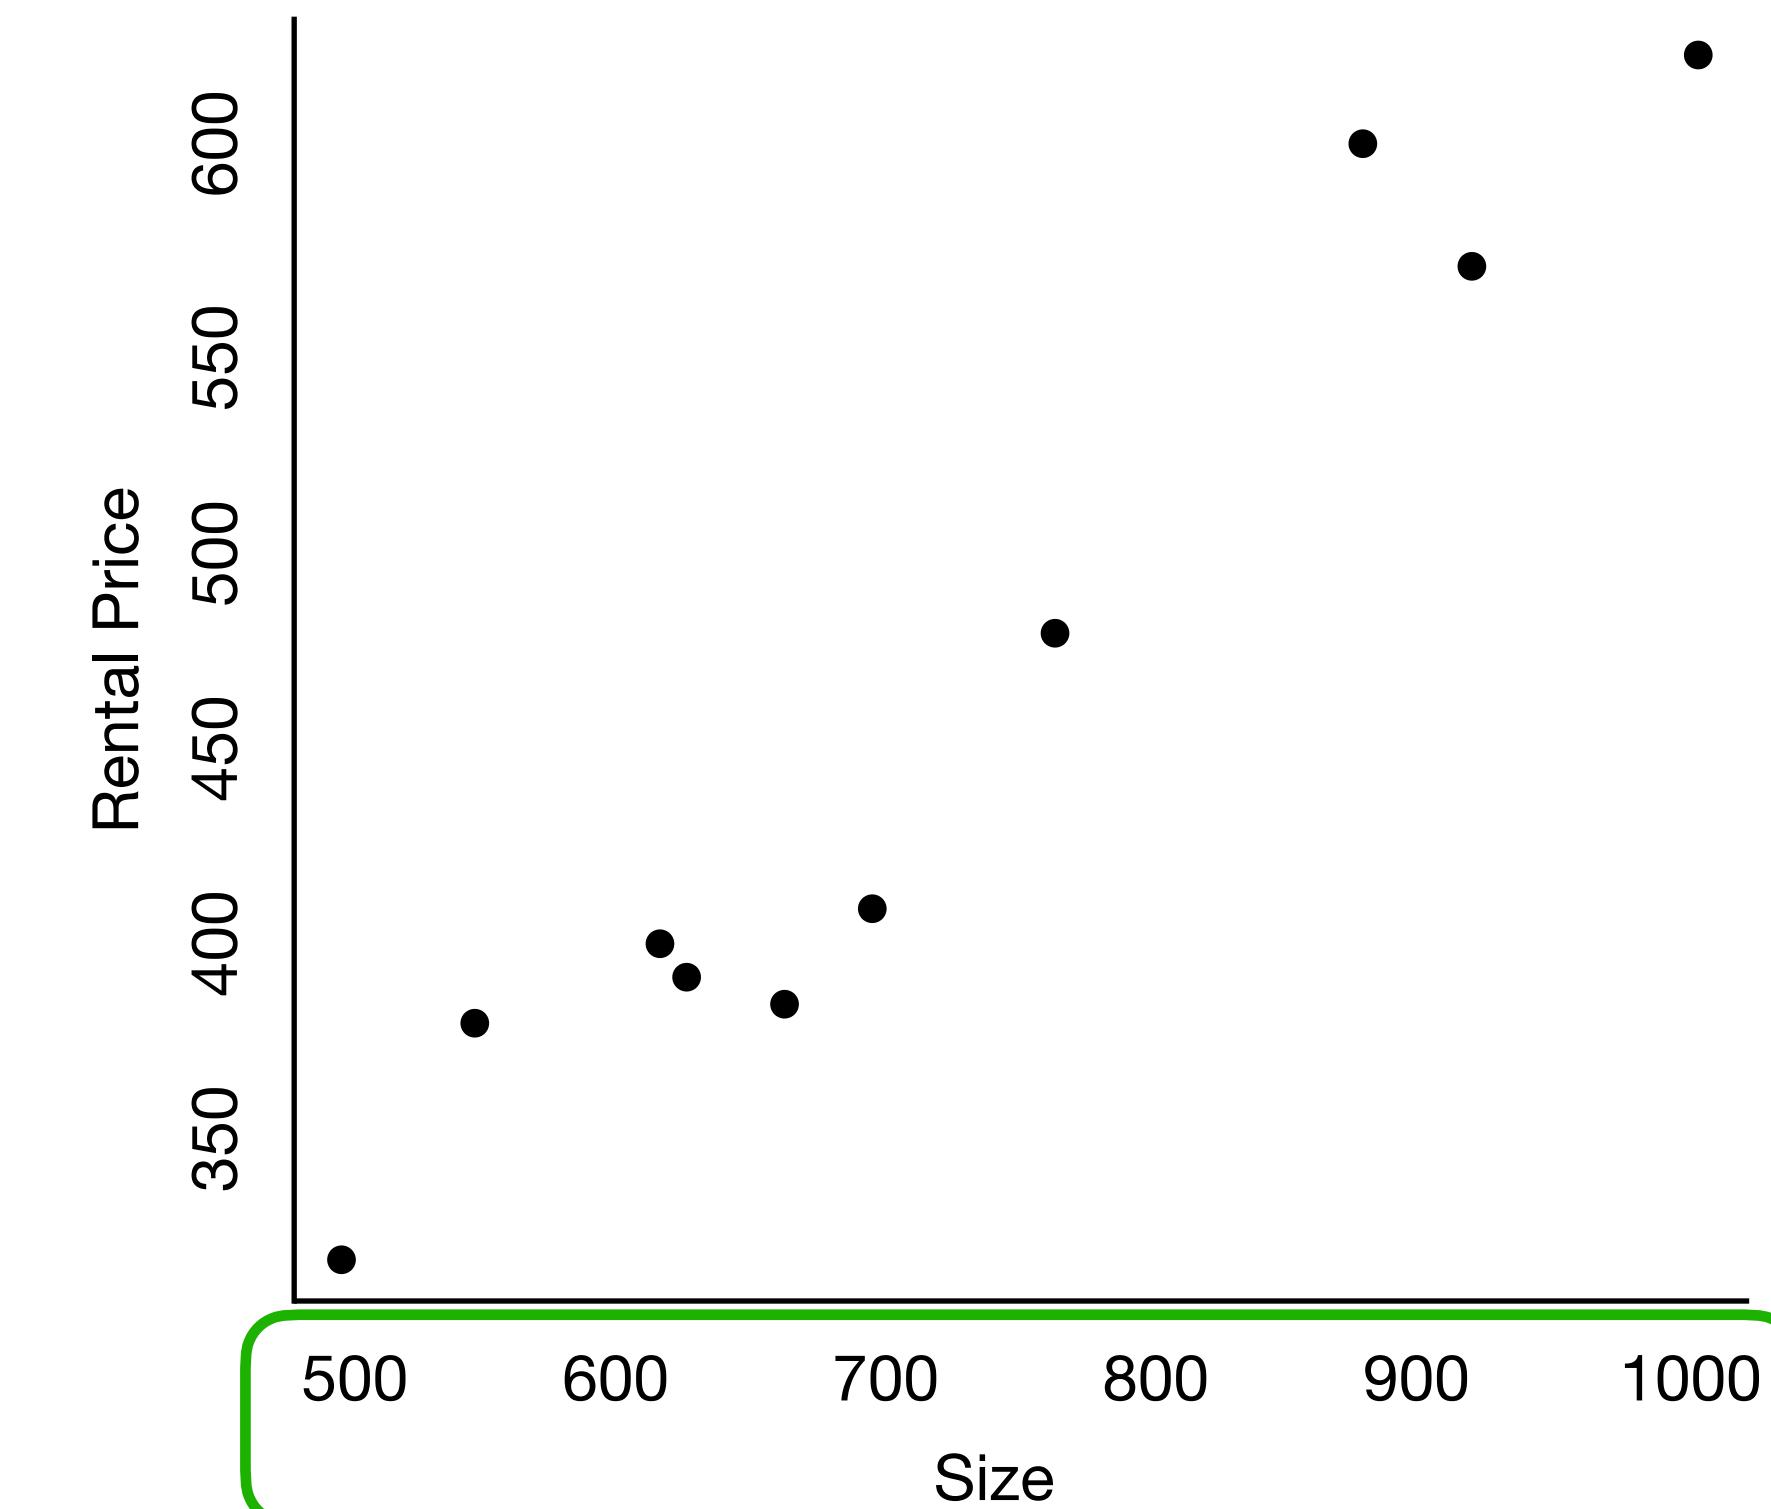
Linear regression. Explained

ID	SIZE	FLOOR	BROADBAND RATE	ENERGY RATING	RENTAL PRICE
1	500	4	8	C	320
2	550	7	50	A	380
3	620	9	7	A	400
4	630	5	24	B	390
5	665	8	100	C	385
6	700	4	8	B	410
7	770	10	7	B	480
8	880	12	50	A	600
9	920	14	8	C	570
10	1,000	9	24	B	620



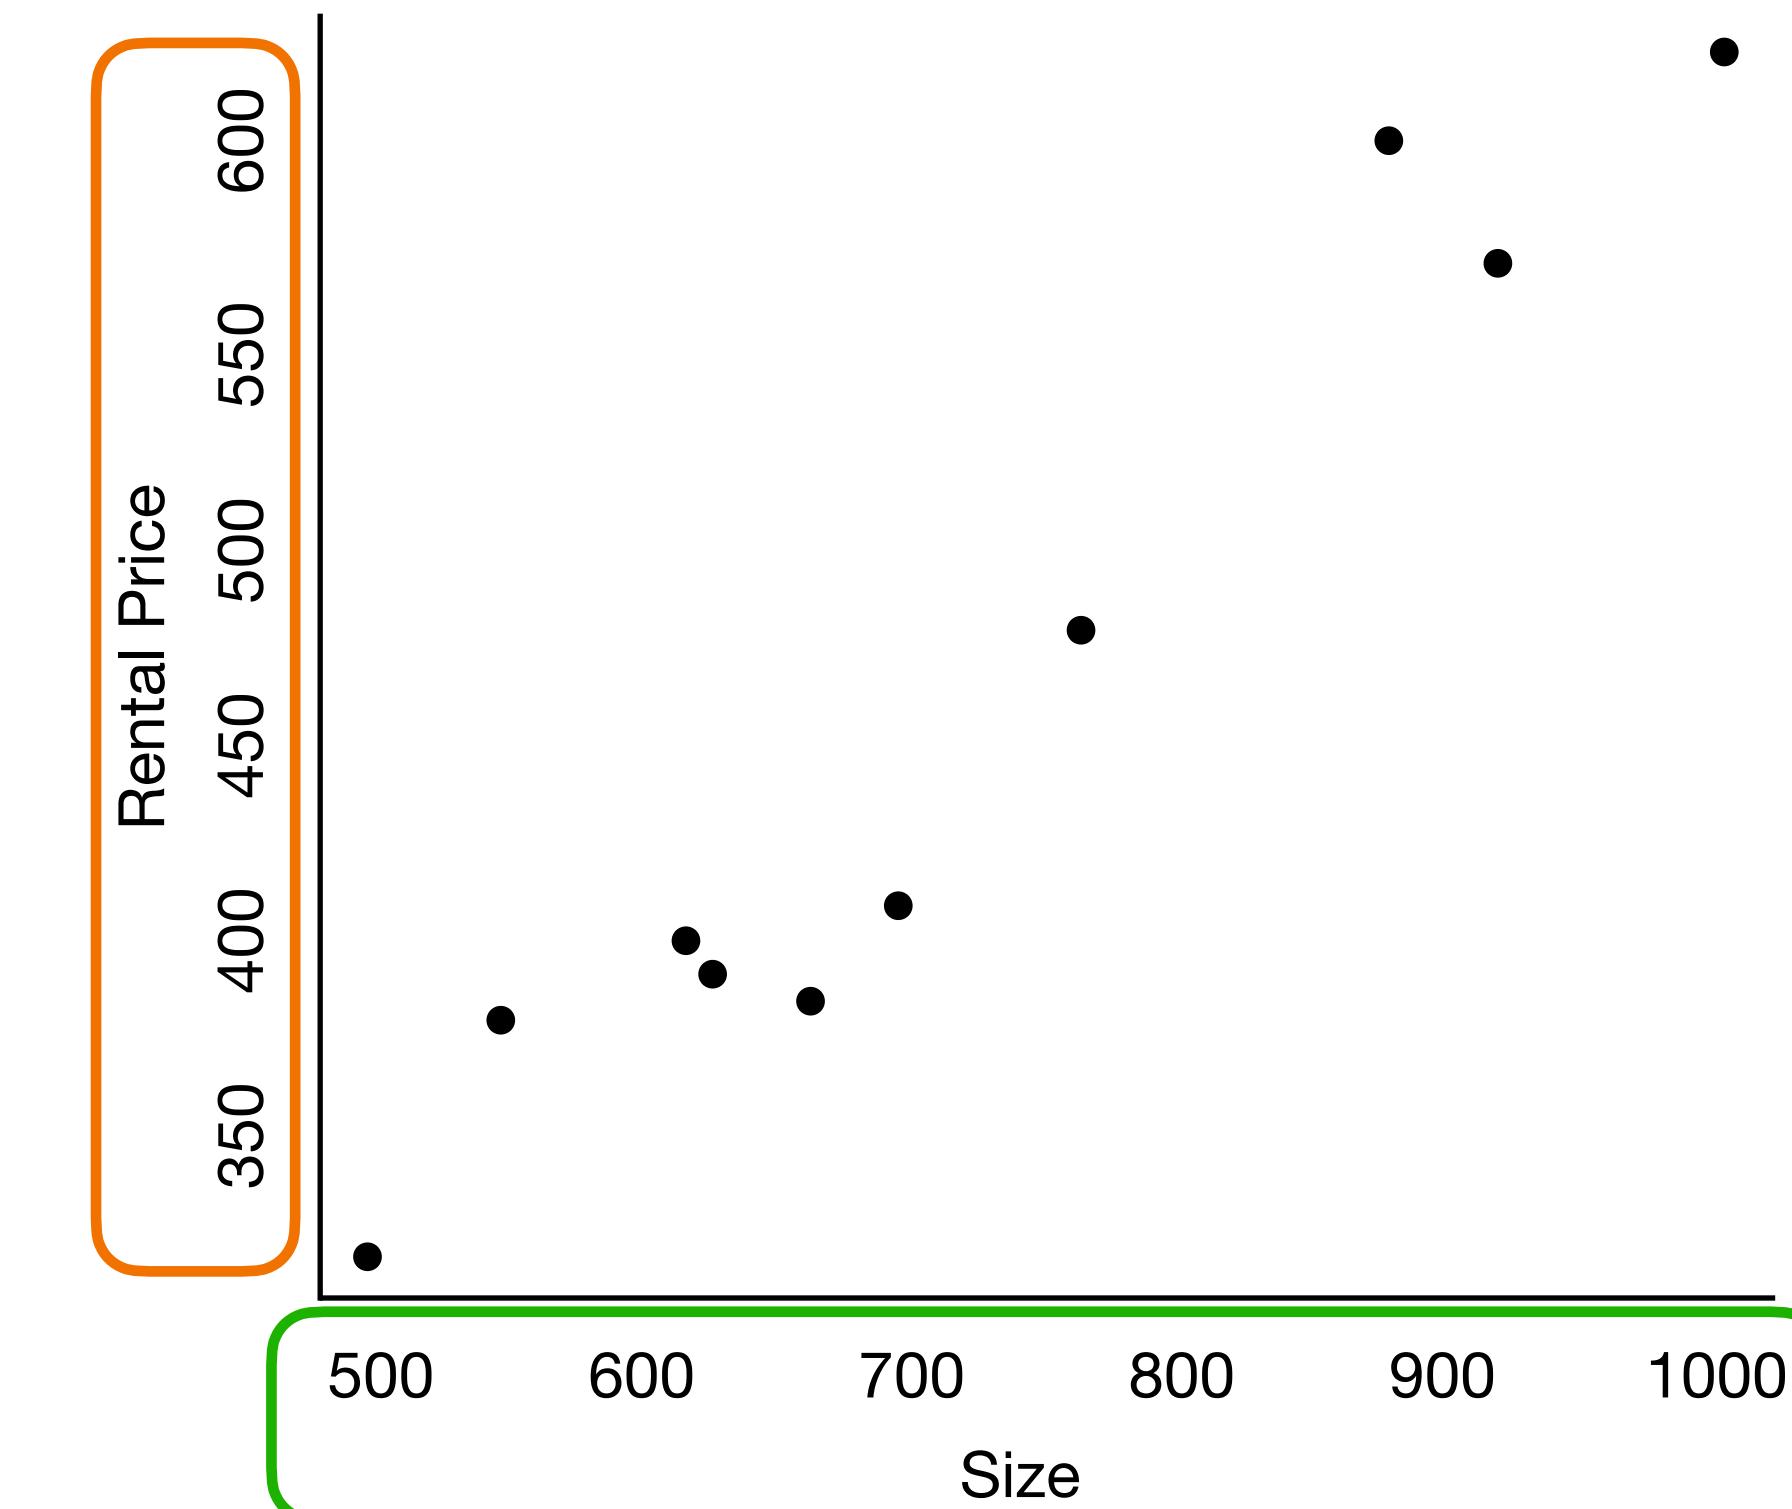
Linear regression. Explained

ID	SIZE	FLOOR	BROADBAND RATE	ENERGY RATING	RENTAL PRICE
1	500	4	8	C	320
2	550	7	50	A	380
3	620	9	7	A	400
4	630	5	24	B	390
5	665	8	100	C	385
6	700	4	8	B	410
7	770	10	7	B	480
8	880	12	50	A	600
9	920	14	8	C	570
10	1,000	9	24	B	620



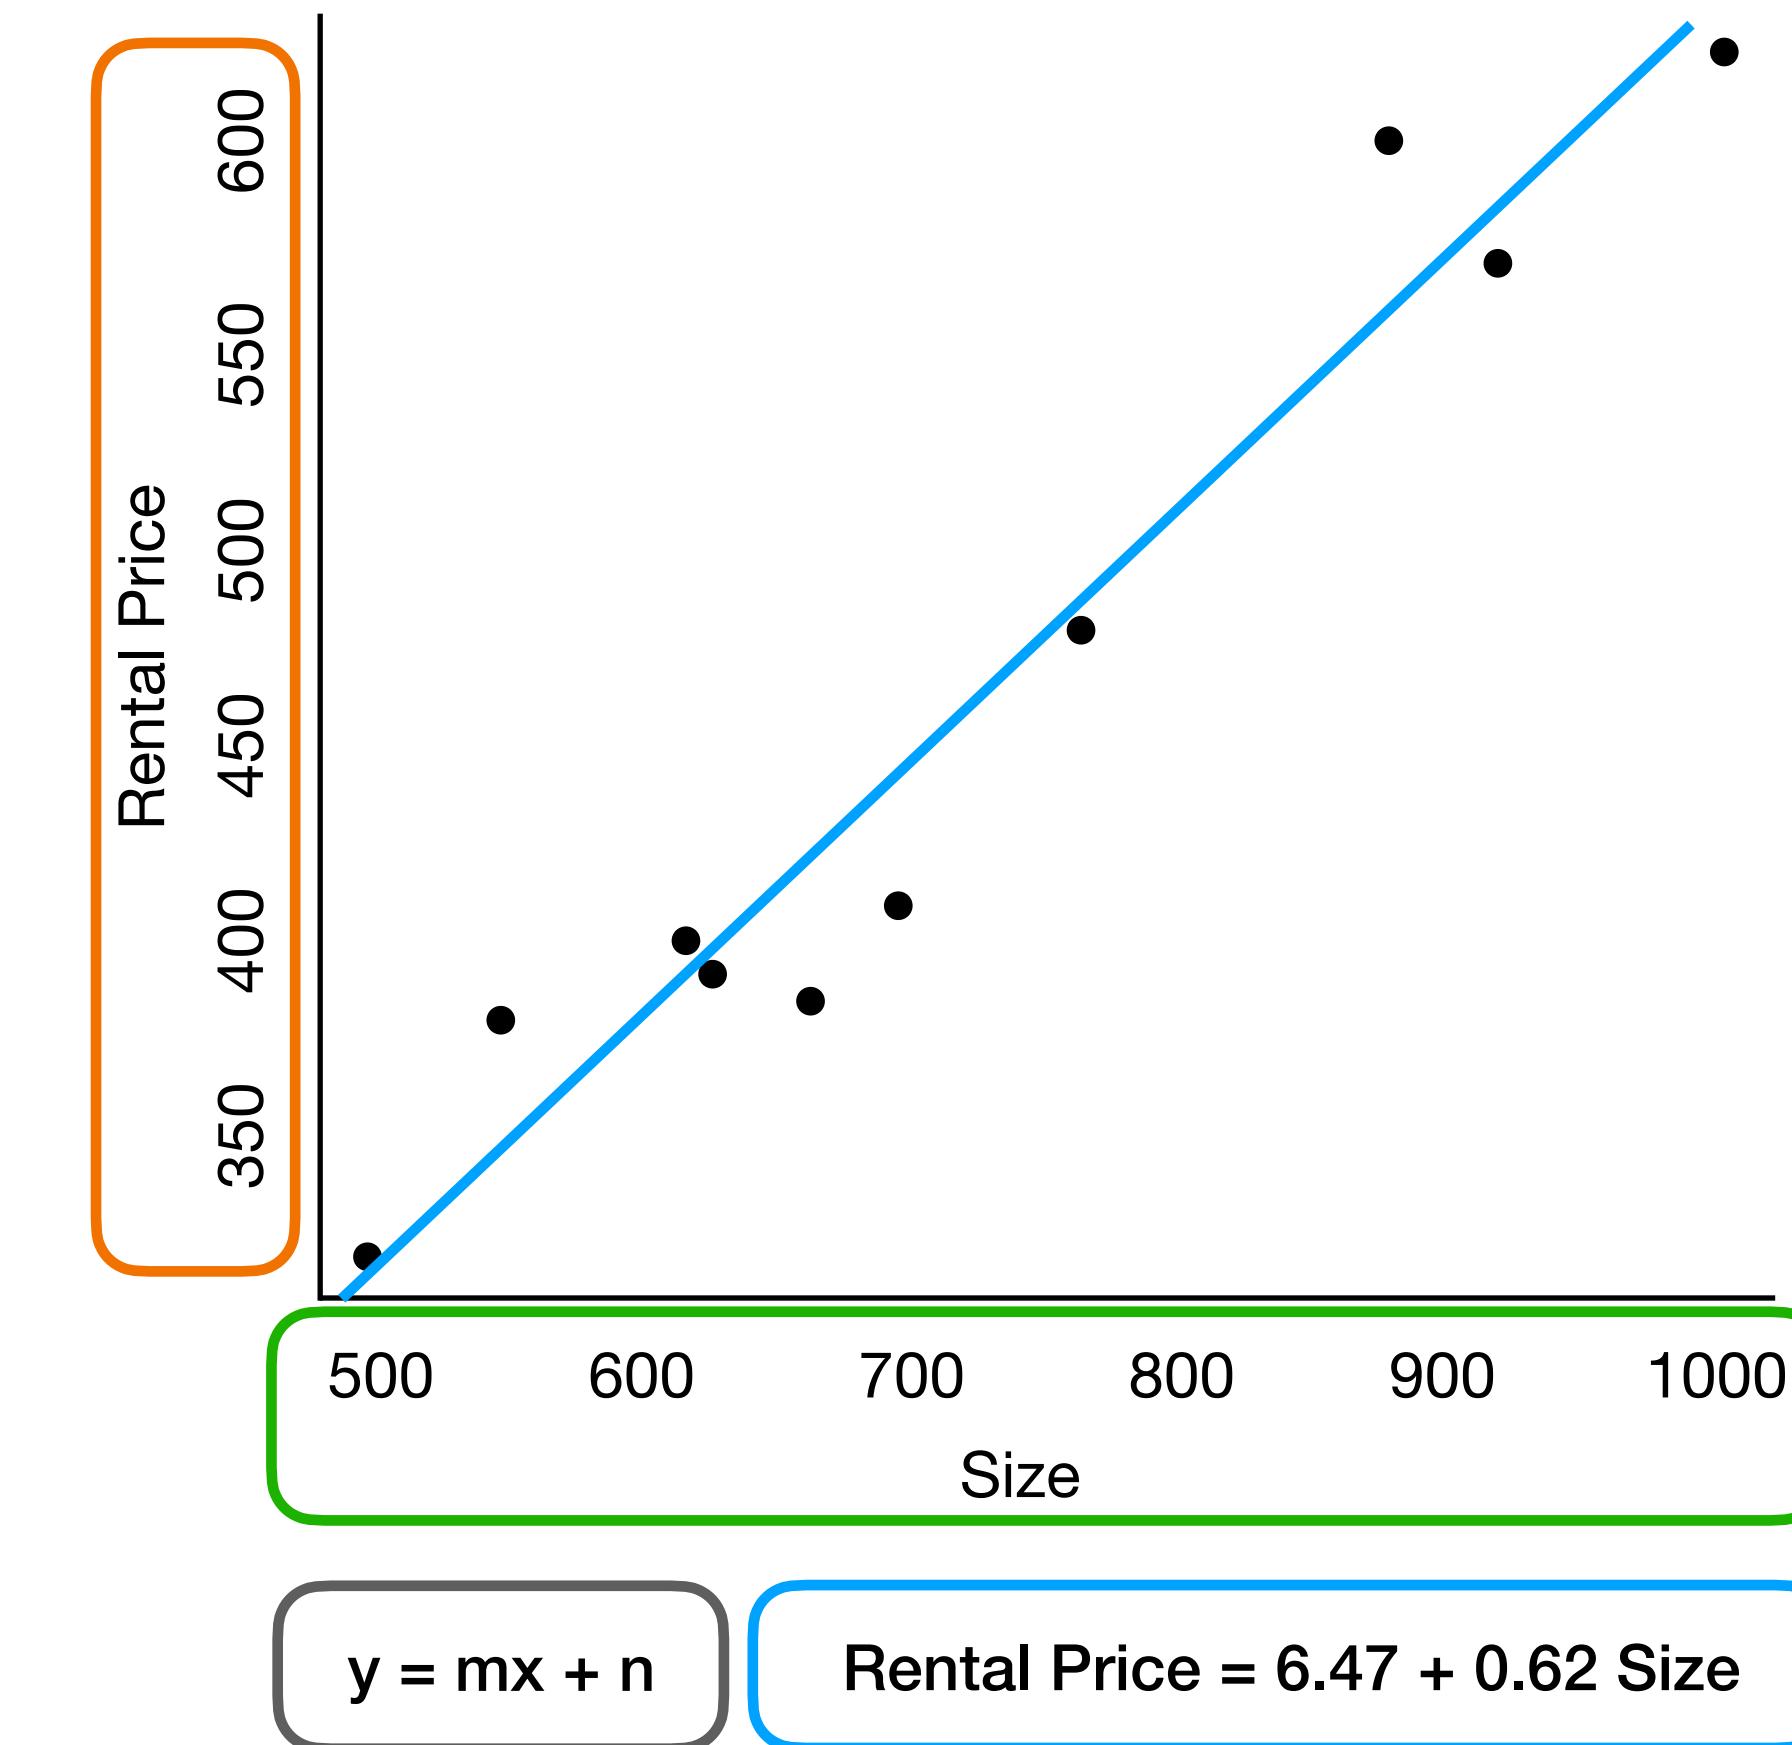
Linear regression. Explained

ID	SIZE	FLOOR	BROADBAND RATE	ENERGY RATING	RENTAL PRICE
1	500	4	8	C	320
2	550	7	50	A	380
3	620	9	7	A	400
4	630	5	24	B	390
5	665	8	100	C	385
6	700	4	8	B	410
7	770	10	7	B	480
8	880	12	50	A	600
9	920	14	8	C	570
10	1,000	9	24	B	620



Linear regression. Explained

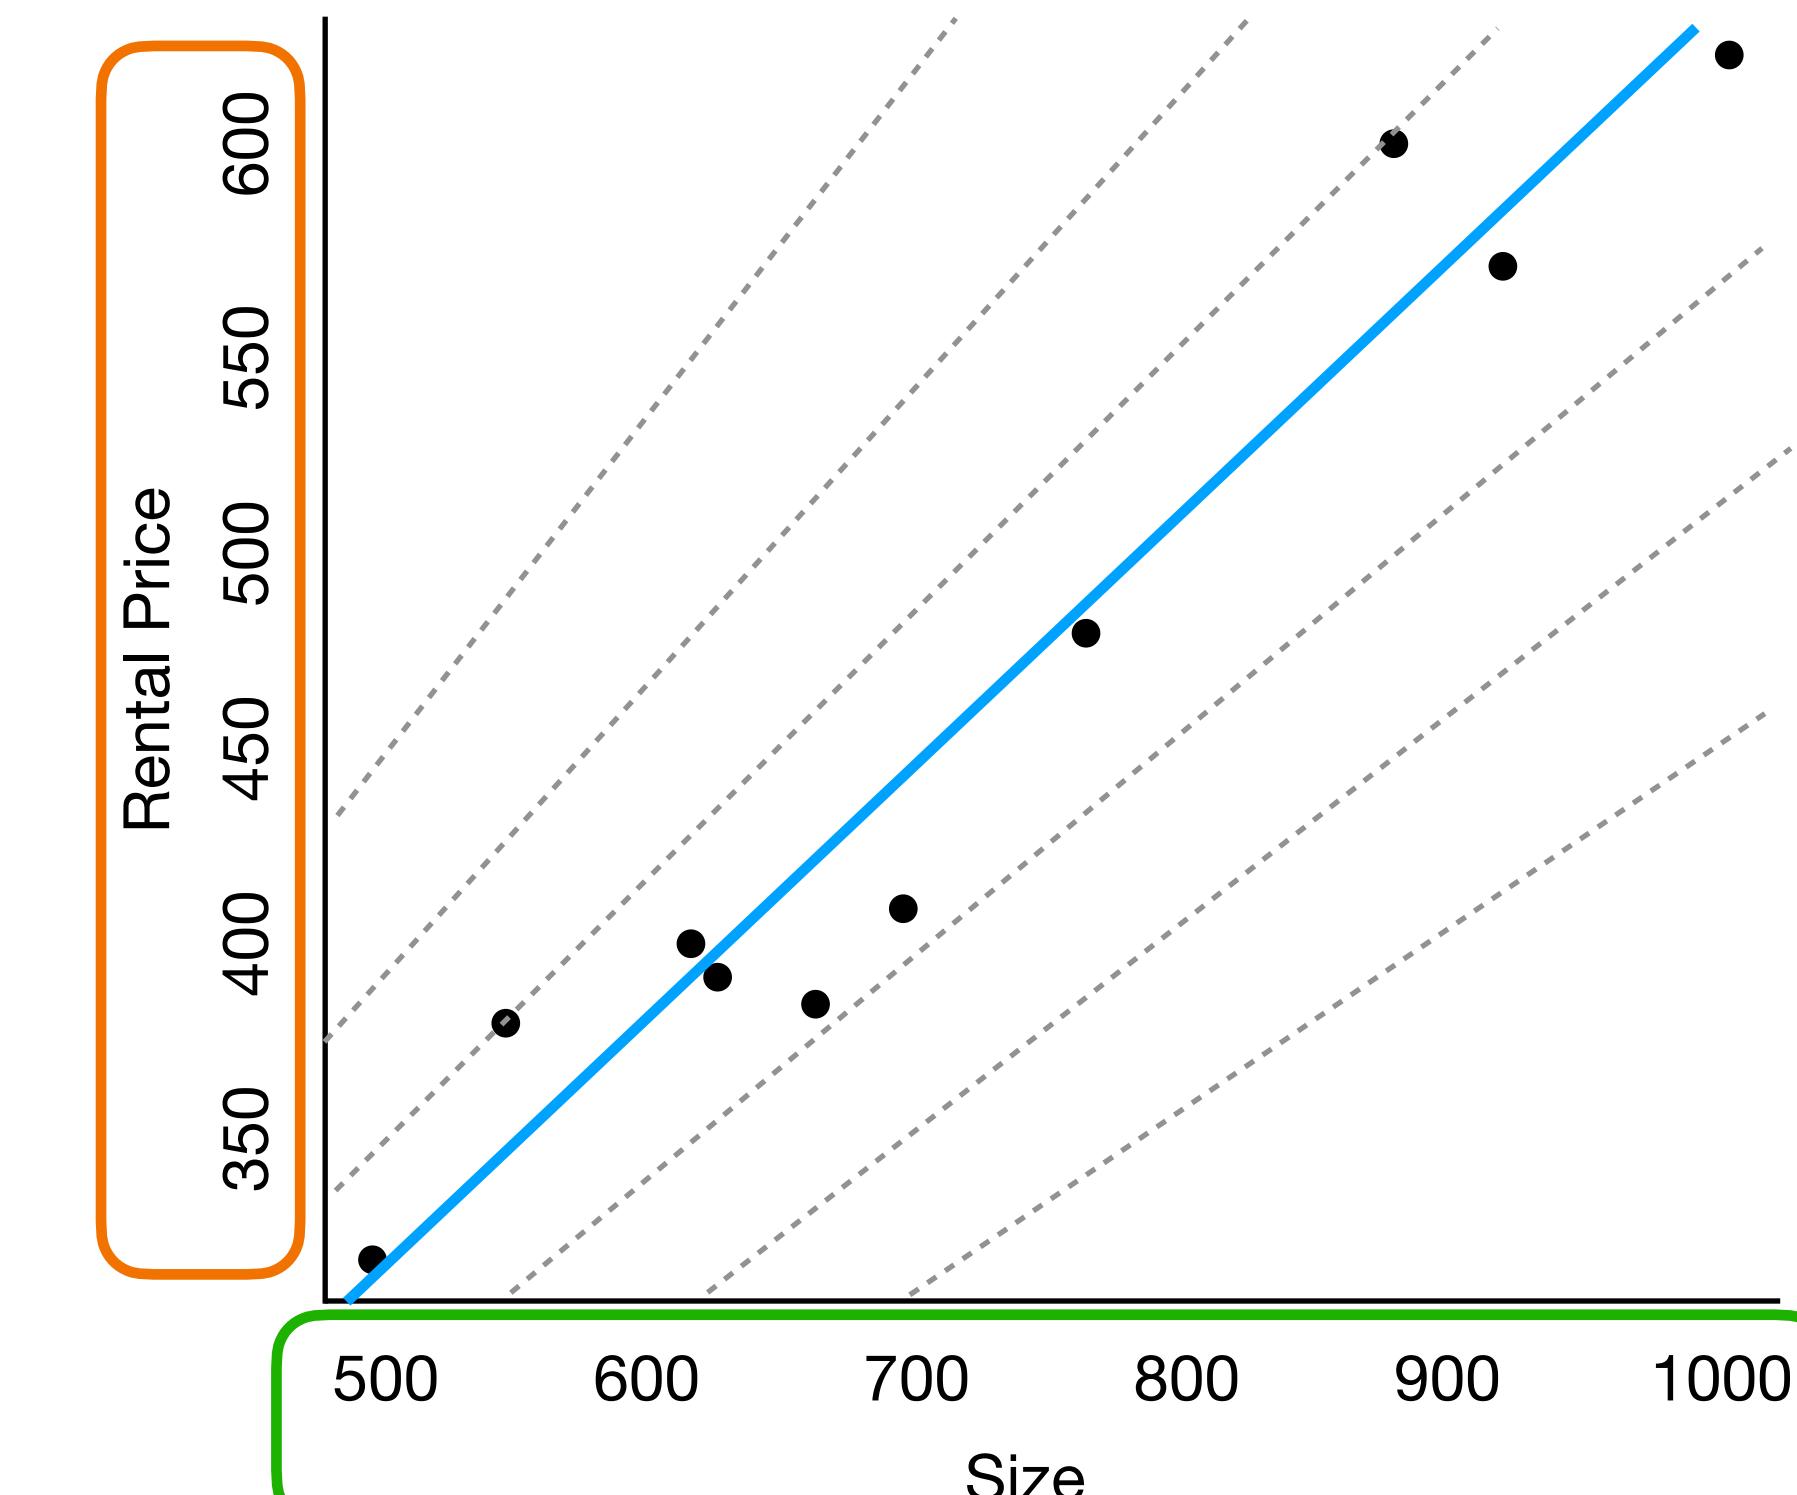
ID	SIZE	FLOOR	BROADBAND RATE	ENERGY RATING	RENTAL PRICE
1	500	4	8	C	320
2	550	7	50	A	380
3	620	9	7	A	400
4	630	5	24	B	390
5	665	8	100	C	385
6	700	4	8	B	410
7	770	10	7	B	480
8	880	12	50	A	600
9	920	14	8	C	570
10	1,000	9	24	B	620



Linear regression. Explained

How do we choose a line? Which one fits better?

ID	SIZE	FLOOR	BROADBAND RATE	ENERGY RATING	RENTAL PRICE
1	500	4	8	C	320
2	550	7	50	A	380
3	620	9	7	A	400
4	630	5	24	B	390
5	665	8	100	C	385
6	700	4	8	B	410
7	770	10	7	B	480
8	880	12	50	A	600
9	920	14	8	C	570
10	1,000	9	24	B	620



$$y = mx + n$$

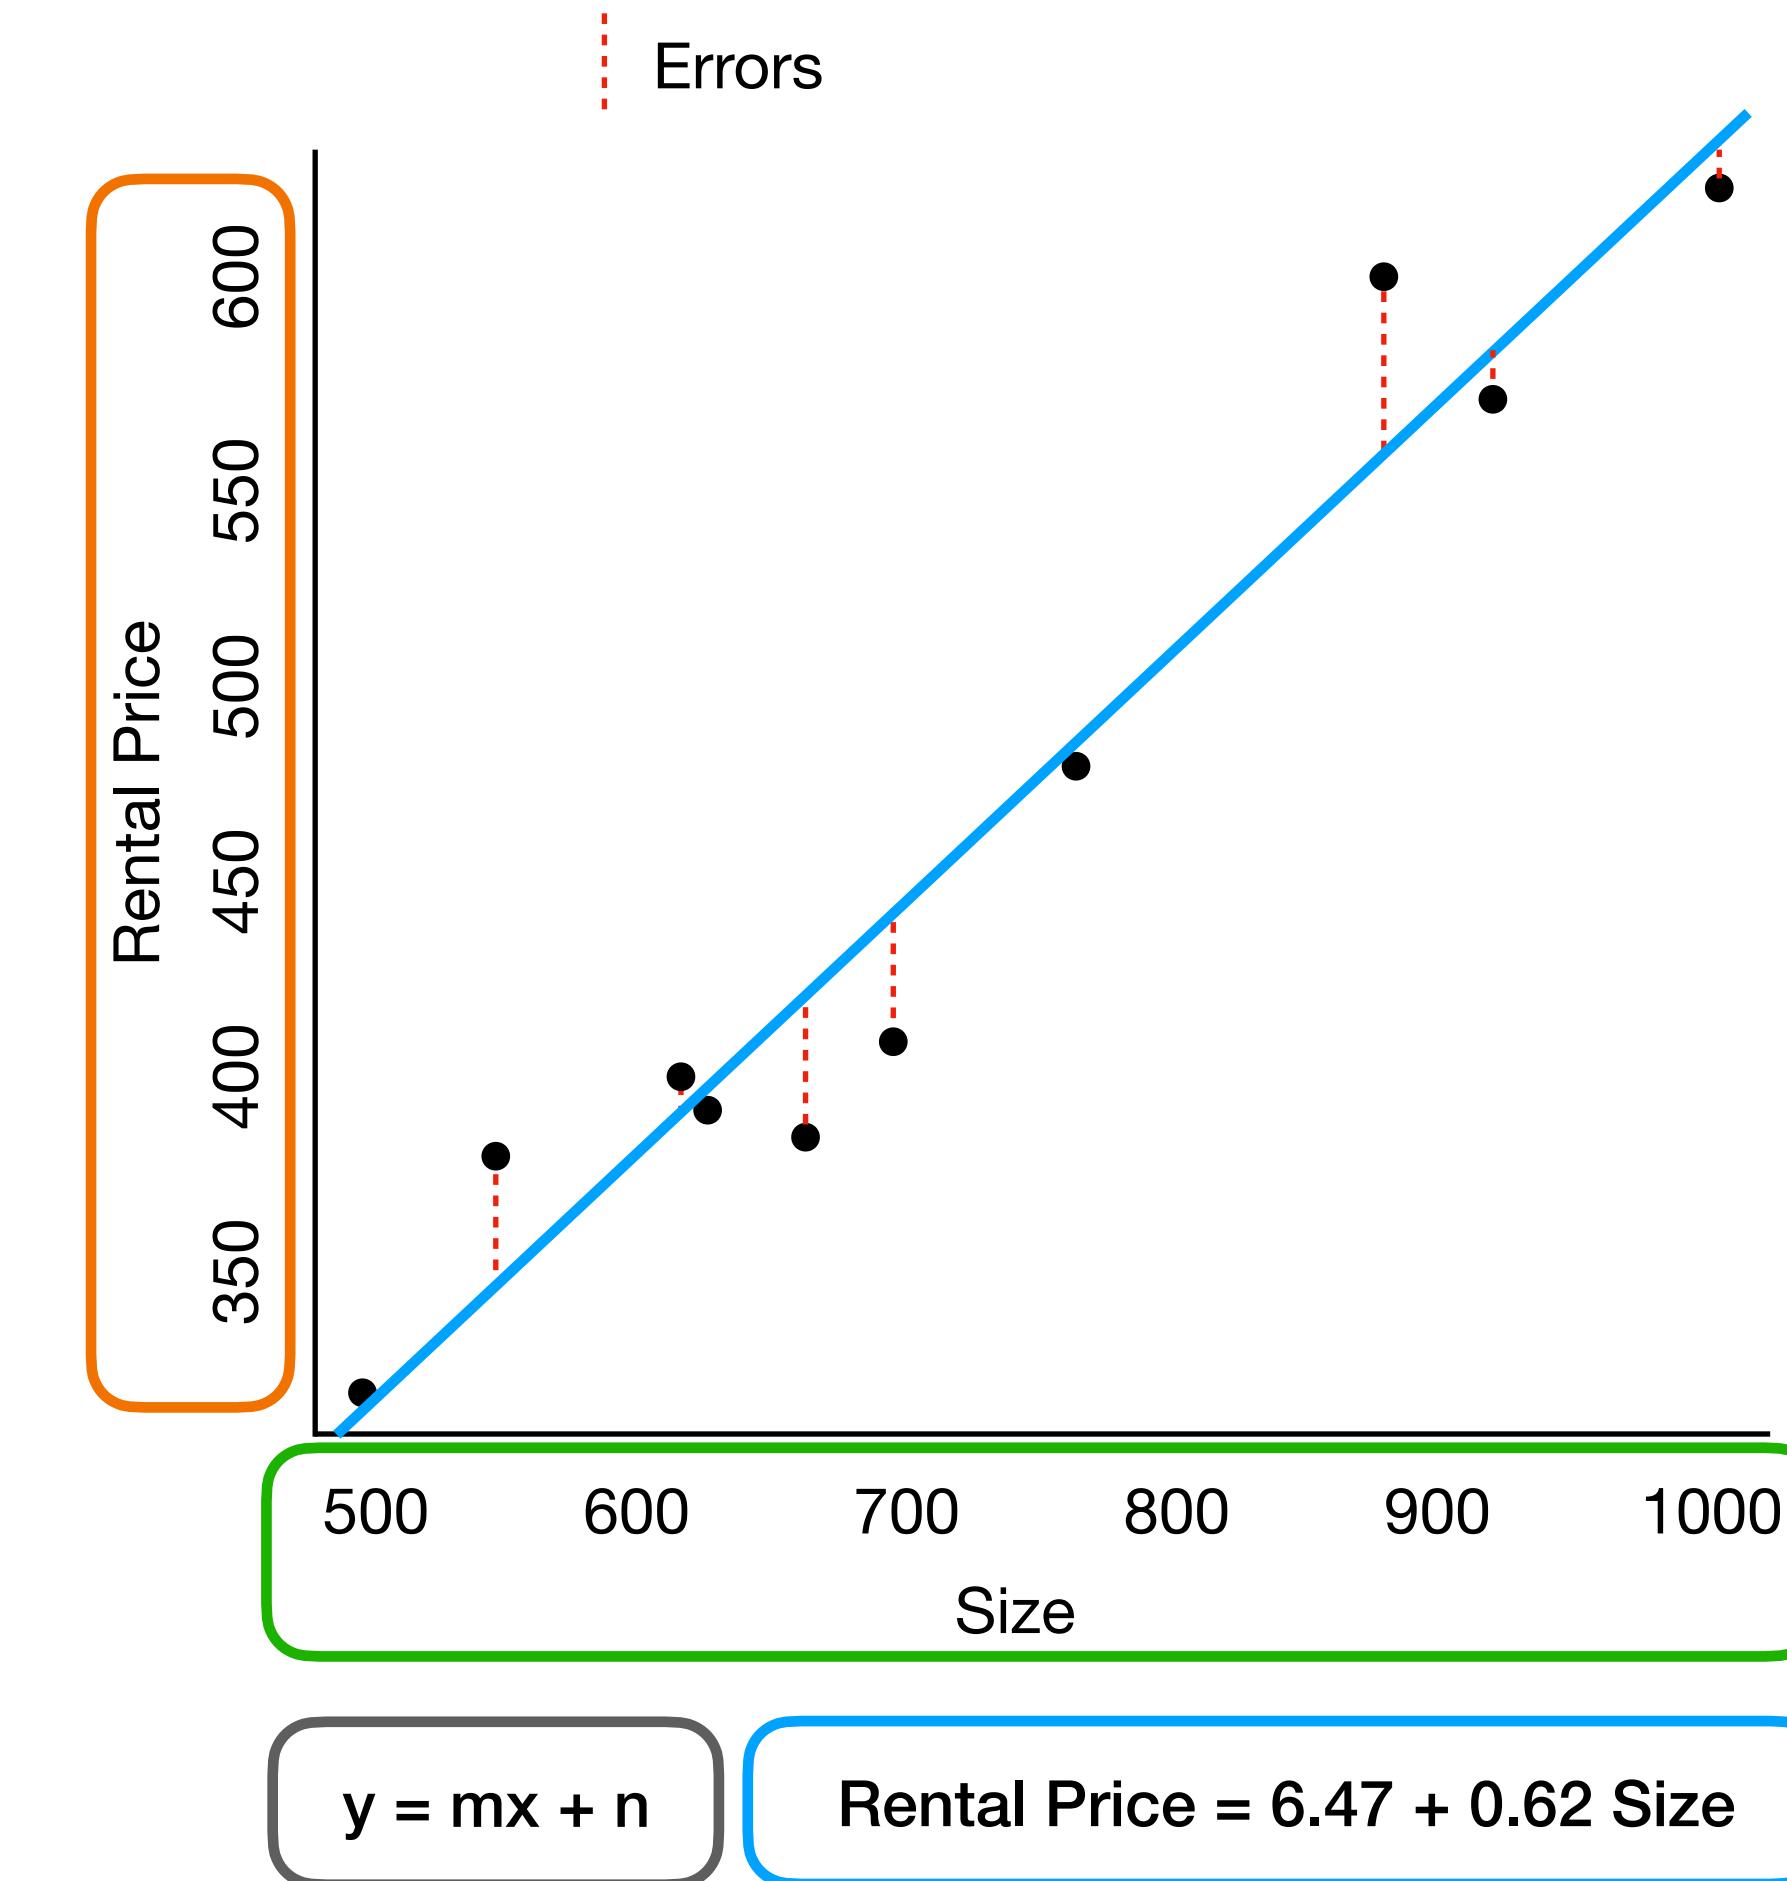
$$\text{Rental Price} = 6.47 + 0.62 \text{ Size}$$

Exercise

- Do a figure with a scatter plot of Math scores and Writing scores

Measuring the error. Sum of Squared Errors

We need an **error function** that captures the errors between the predictions of the model (blue line) and the actual values in the dataset (Real rental price)

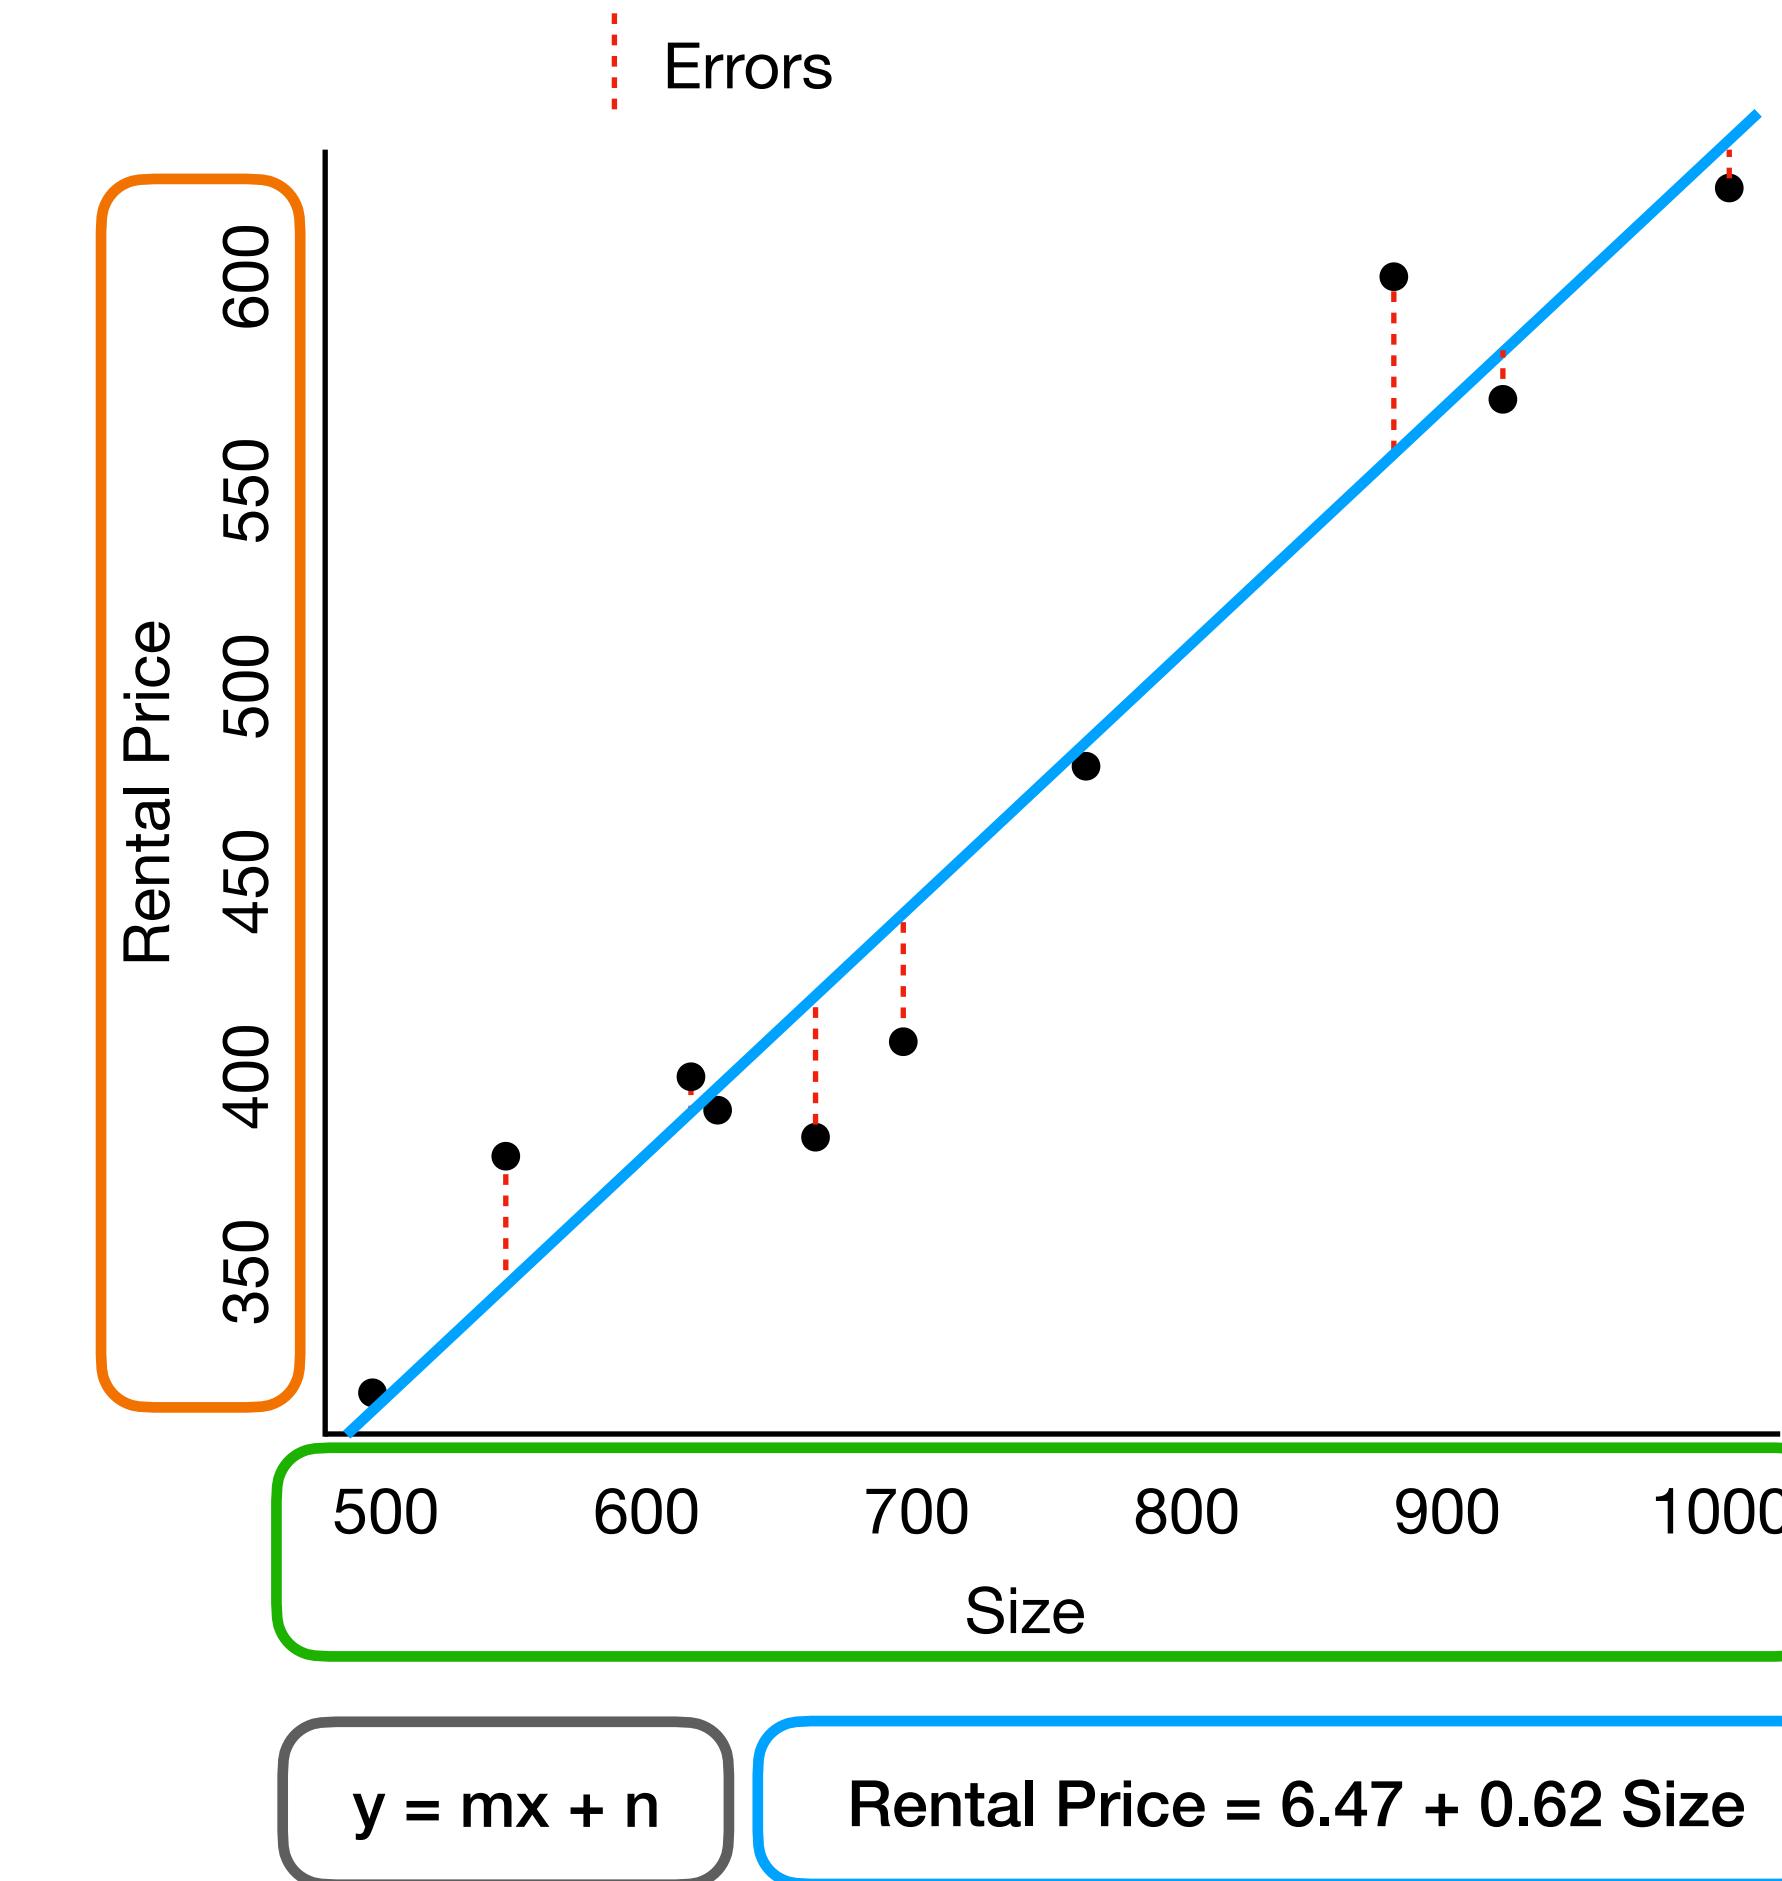


Measuring the error. Sum of Squared Errors

We need an **error function** that captures the errors between the predictions of the model (blue line) and the actual values in the dataset (Real rental price)

Sum of Squared Errors

$$L_2 = \frac{1}{2} \sum_{i=1}^n (y - y_{pred})^2$$



Measuring the error. Sum of Squared Errors

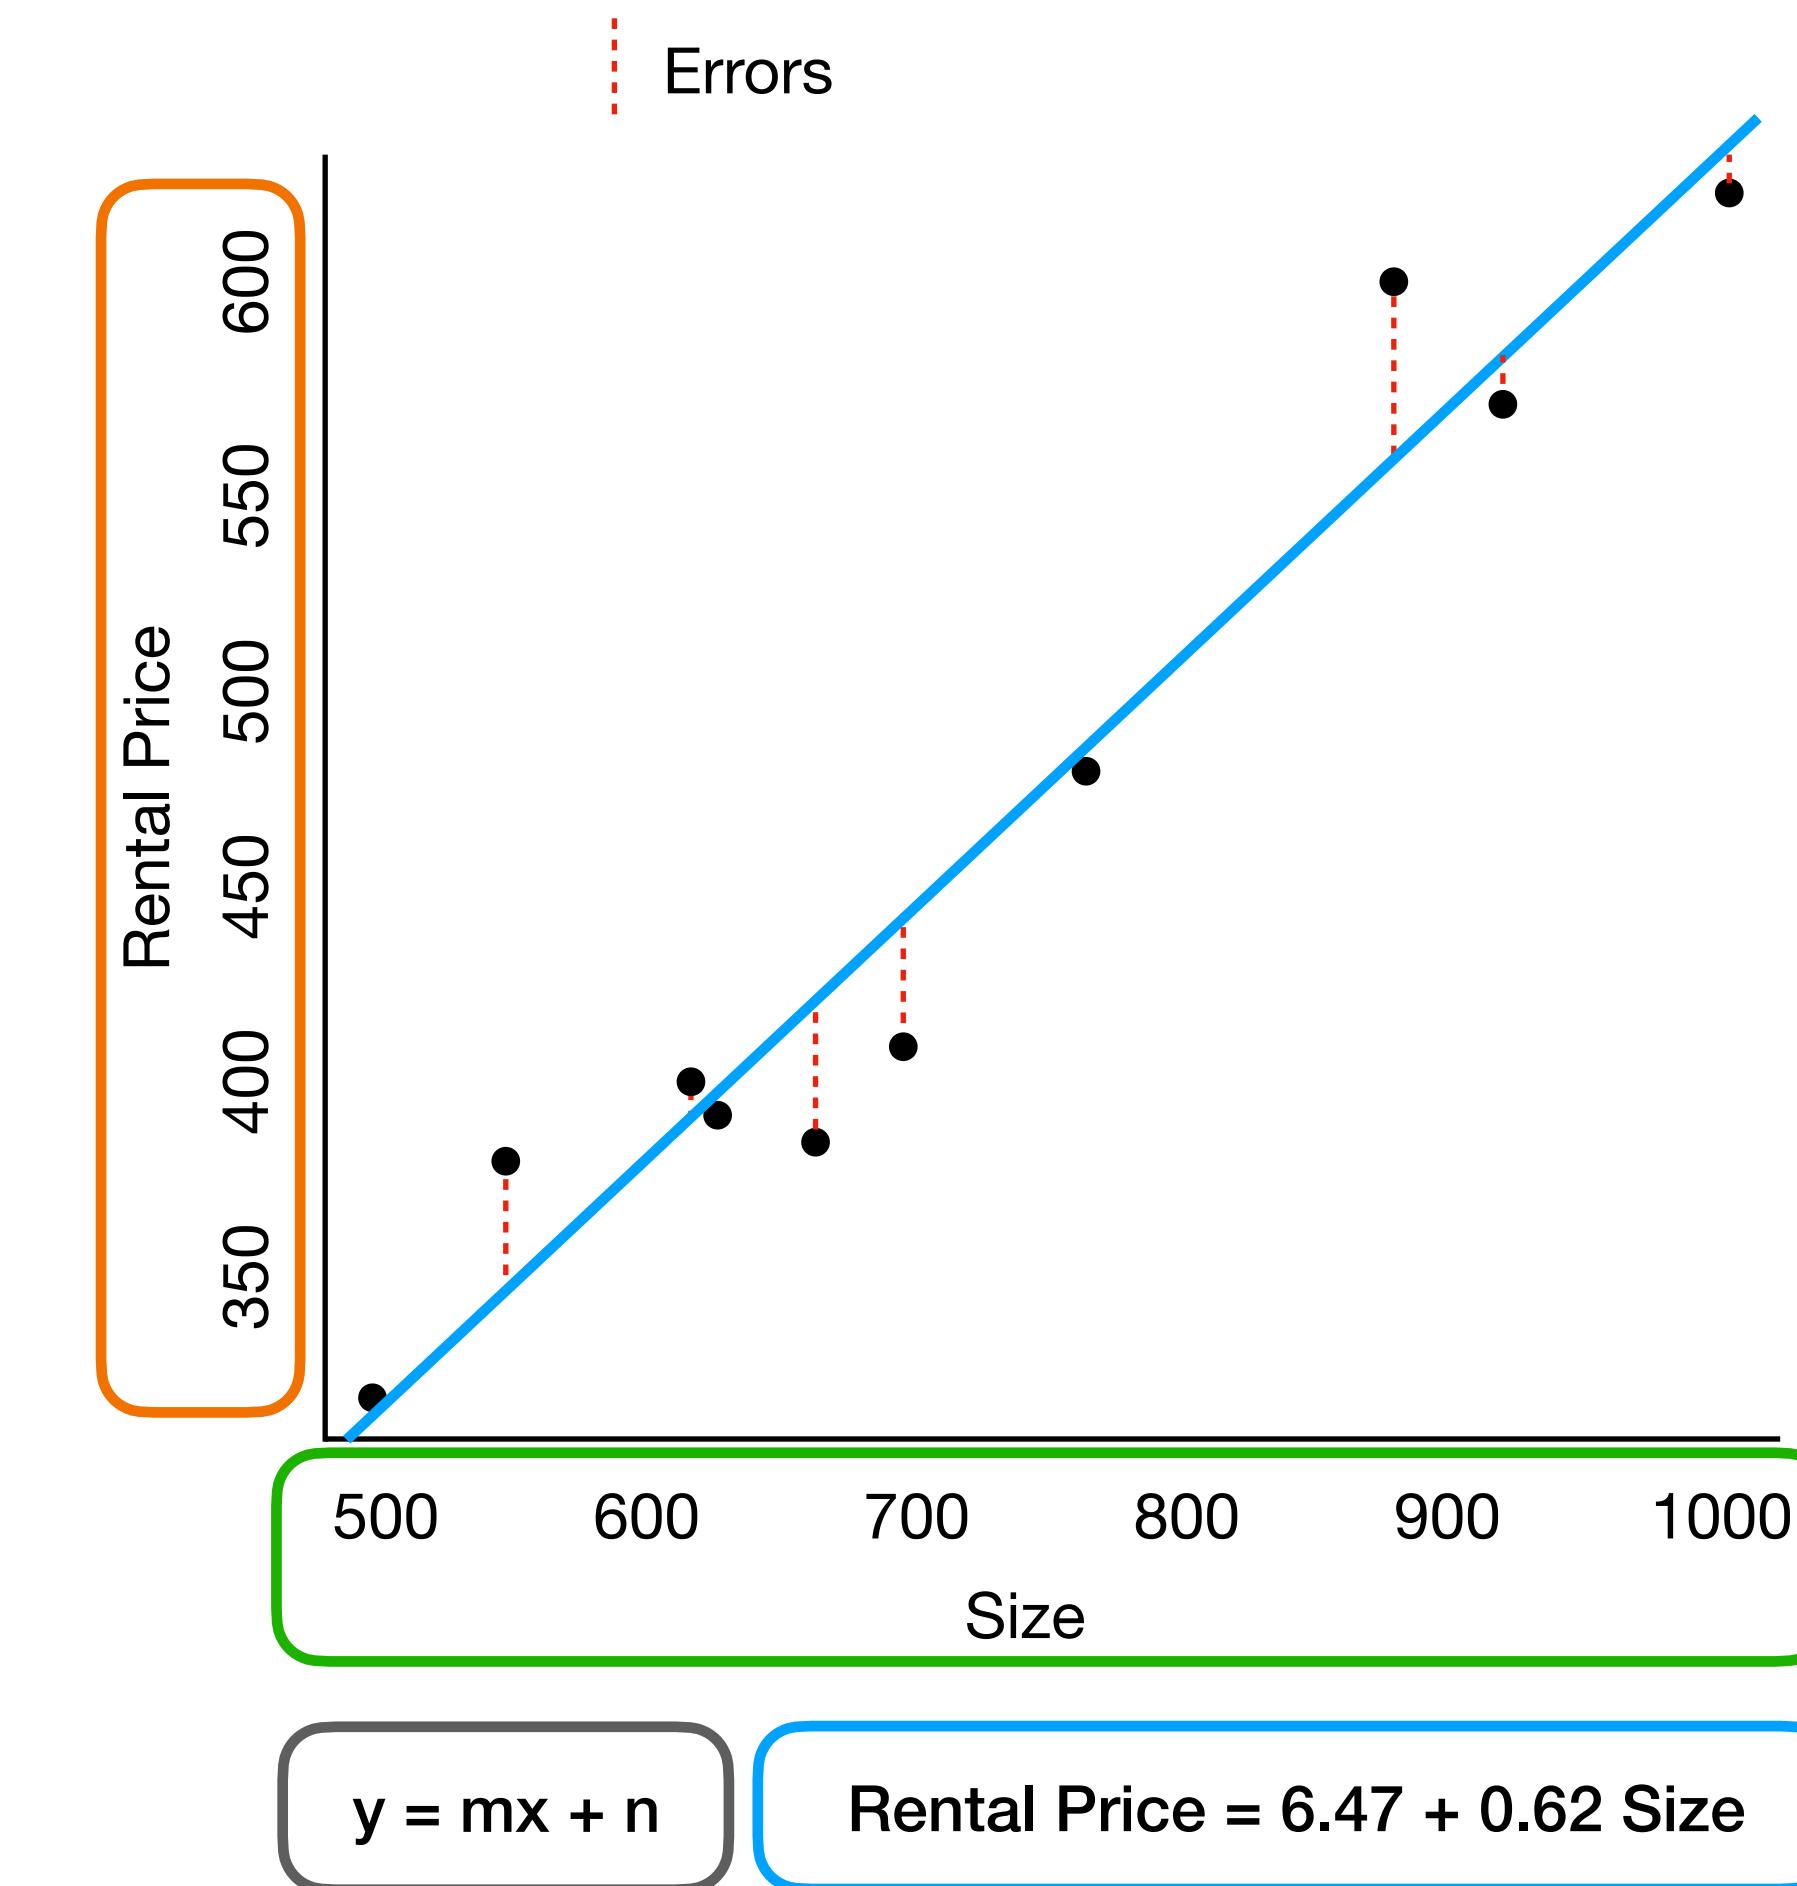
We need an **error function** that captures the errors between the predictions of the model (blue line) and the actual values in the dataset (Real rental price)

Sum of Squared Errors

$$L_2 = \frac{1}{2} \sum_{i=1}^n (y - y_{pred})$$

$$L_2 = \frac{1}{2} \sum_{i=1}^n (y - (mx + n))$$

$$L_2 = \frac{1}{2} \sum_{i=1}^n (y - (0.62 \cdot Size + 6.47))$$

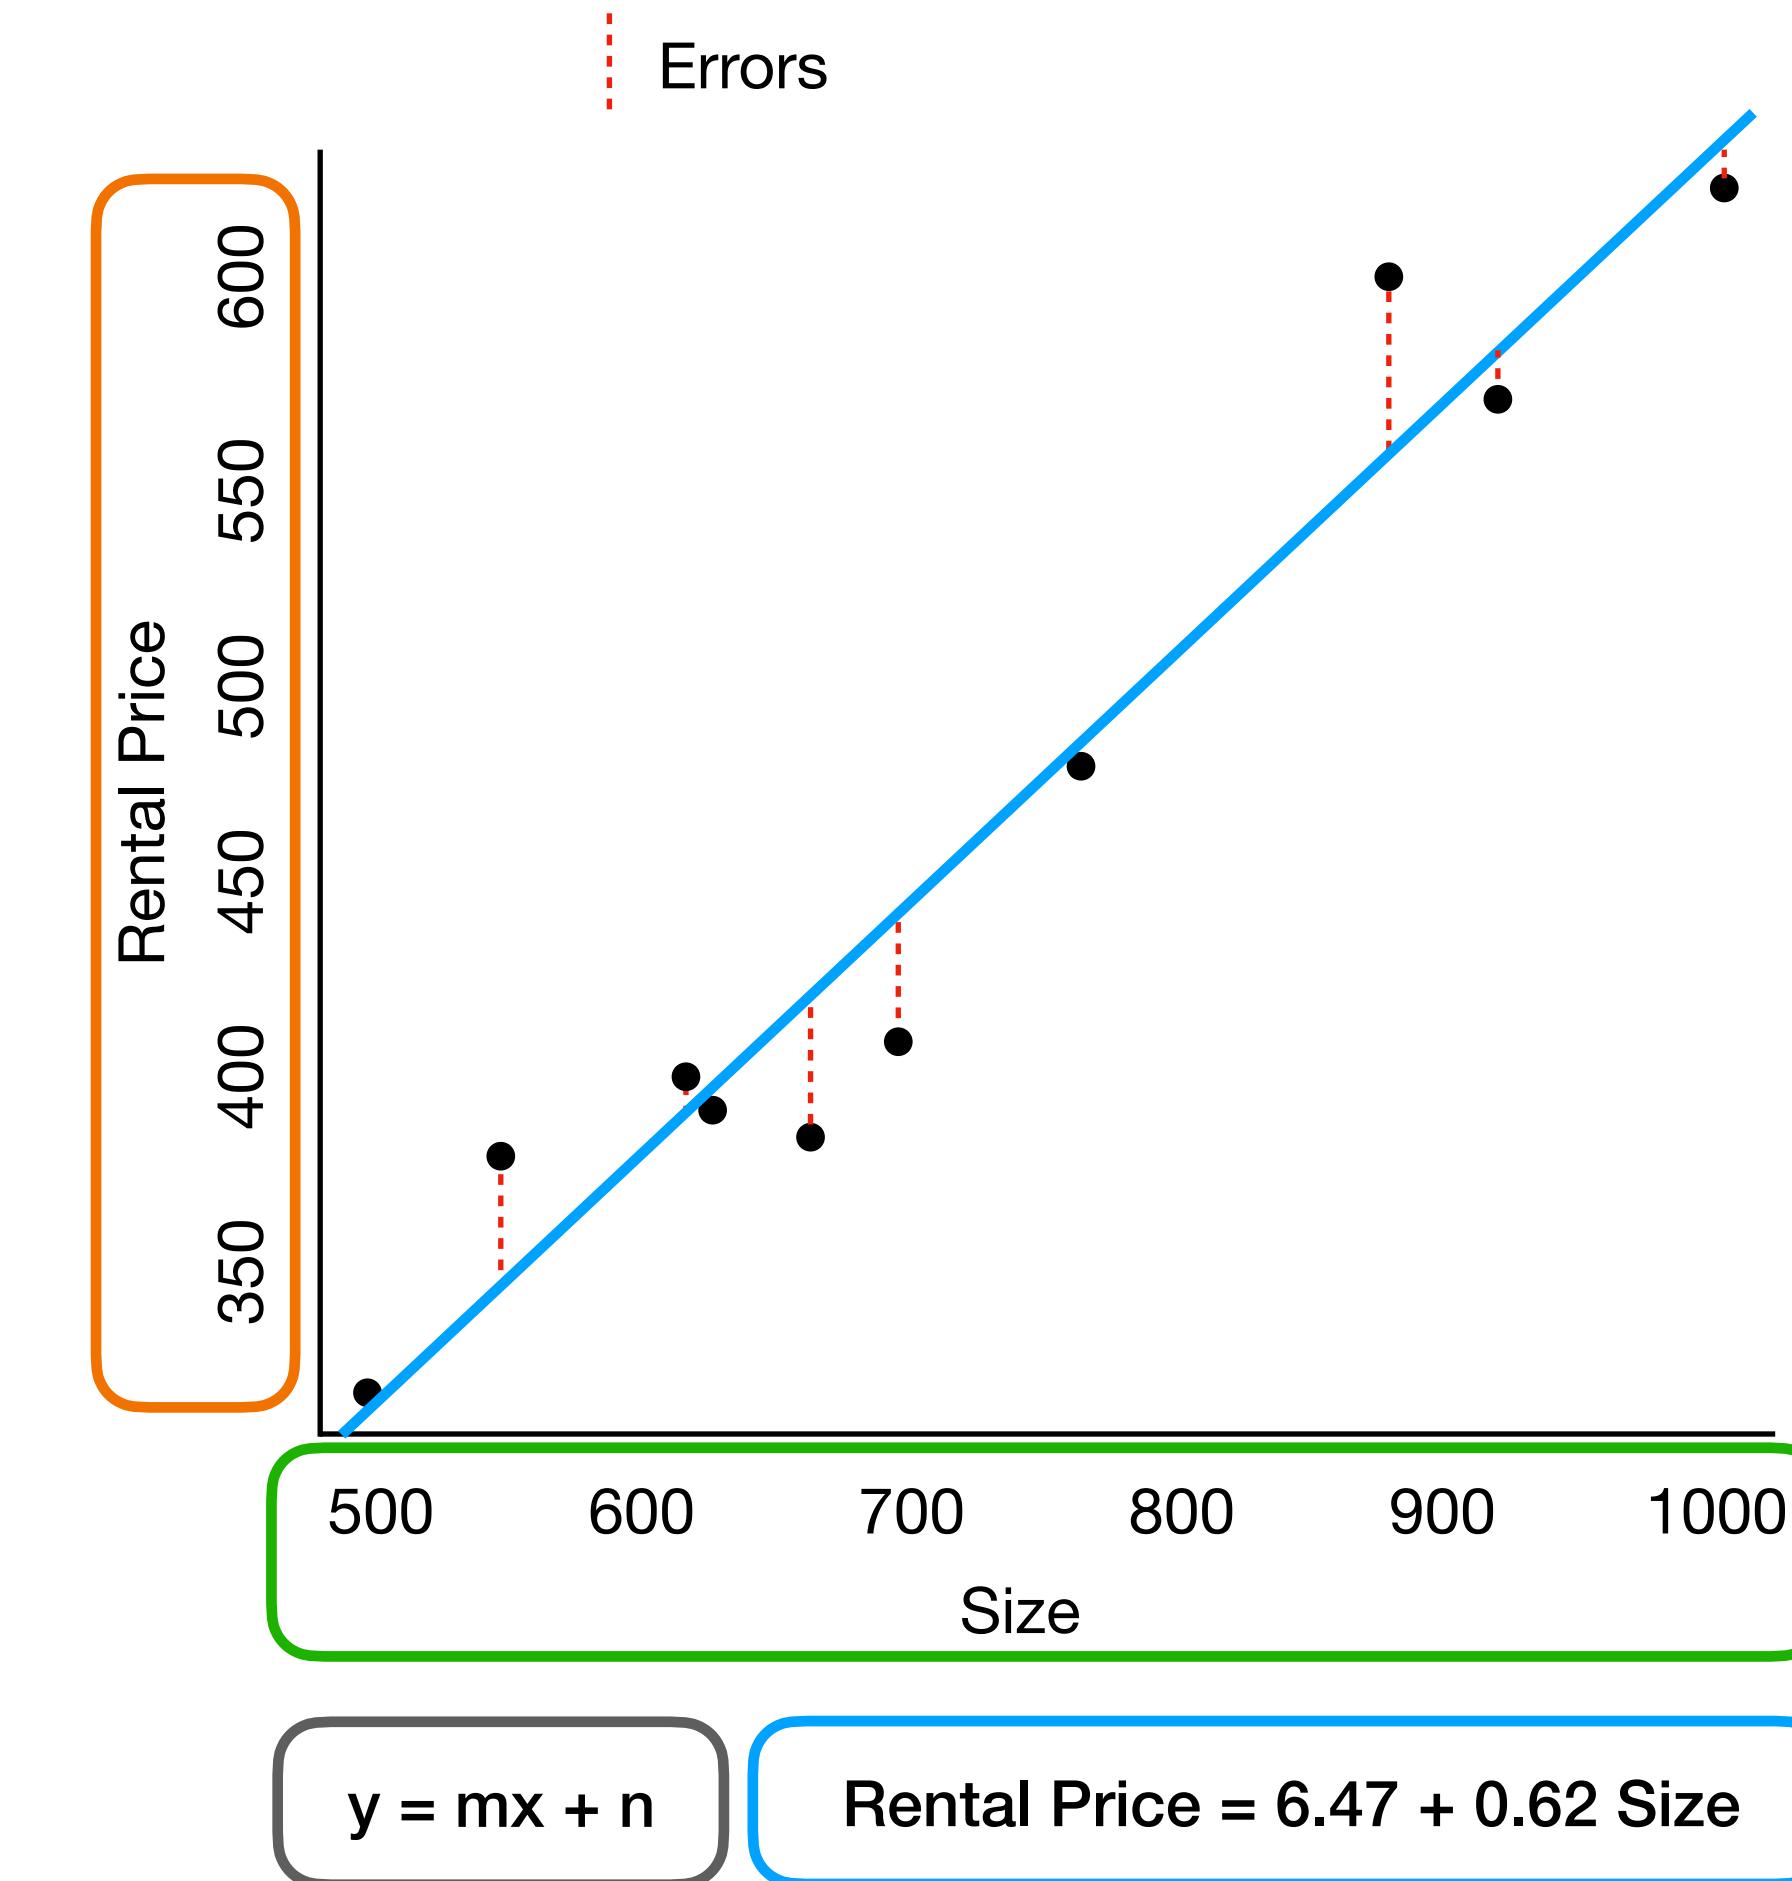


Objective: Find the best fit of a line to the data

We want to obtain the line that minimizes the sum of squared errors.

$$\min(L_2)$$

To find the best parameters for the line that we are searching for (m and n in this case)



Objective: Find the best fit of a line to the data

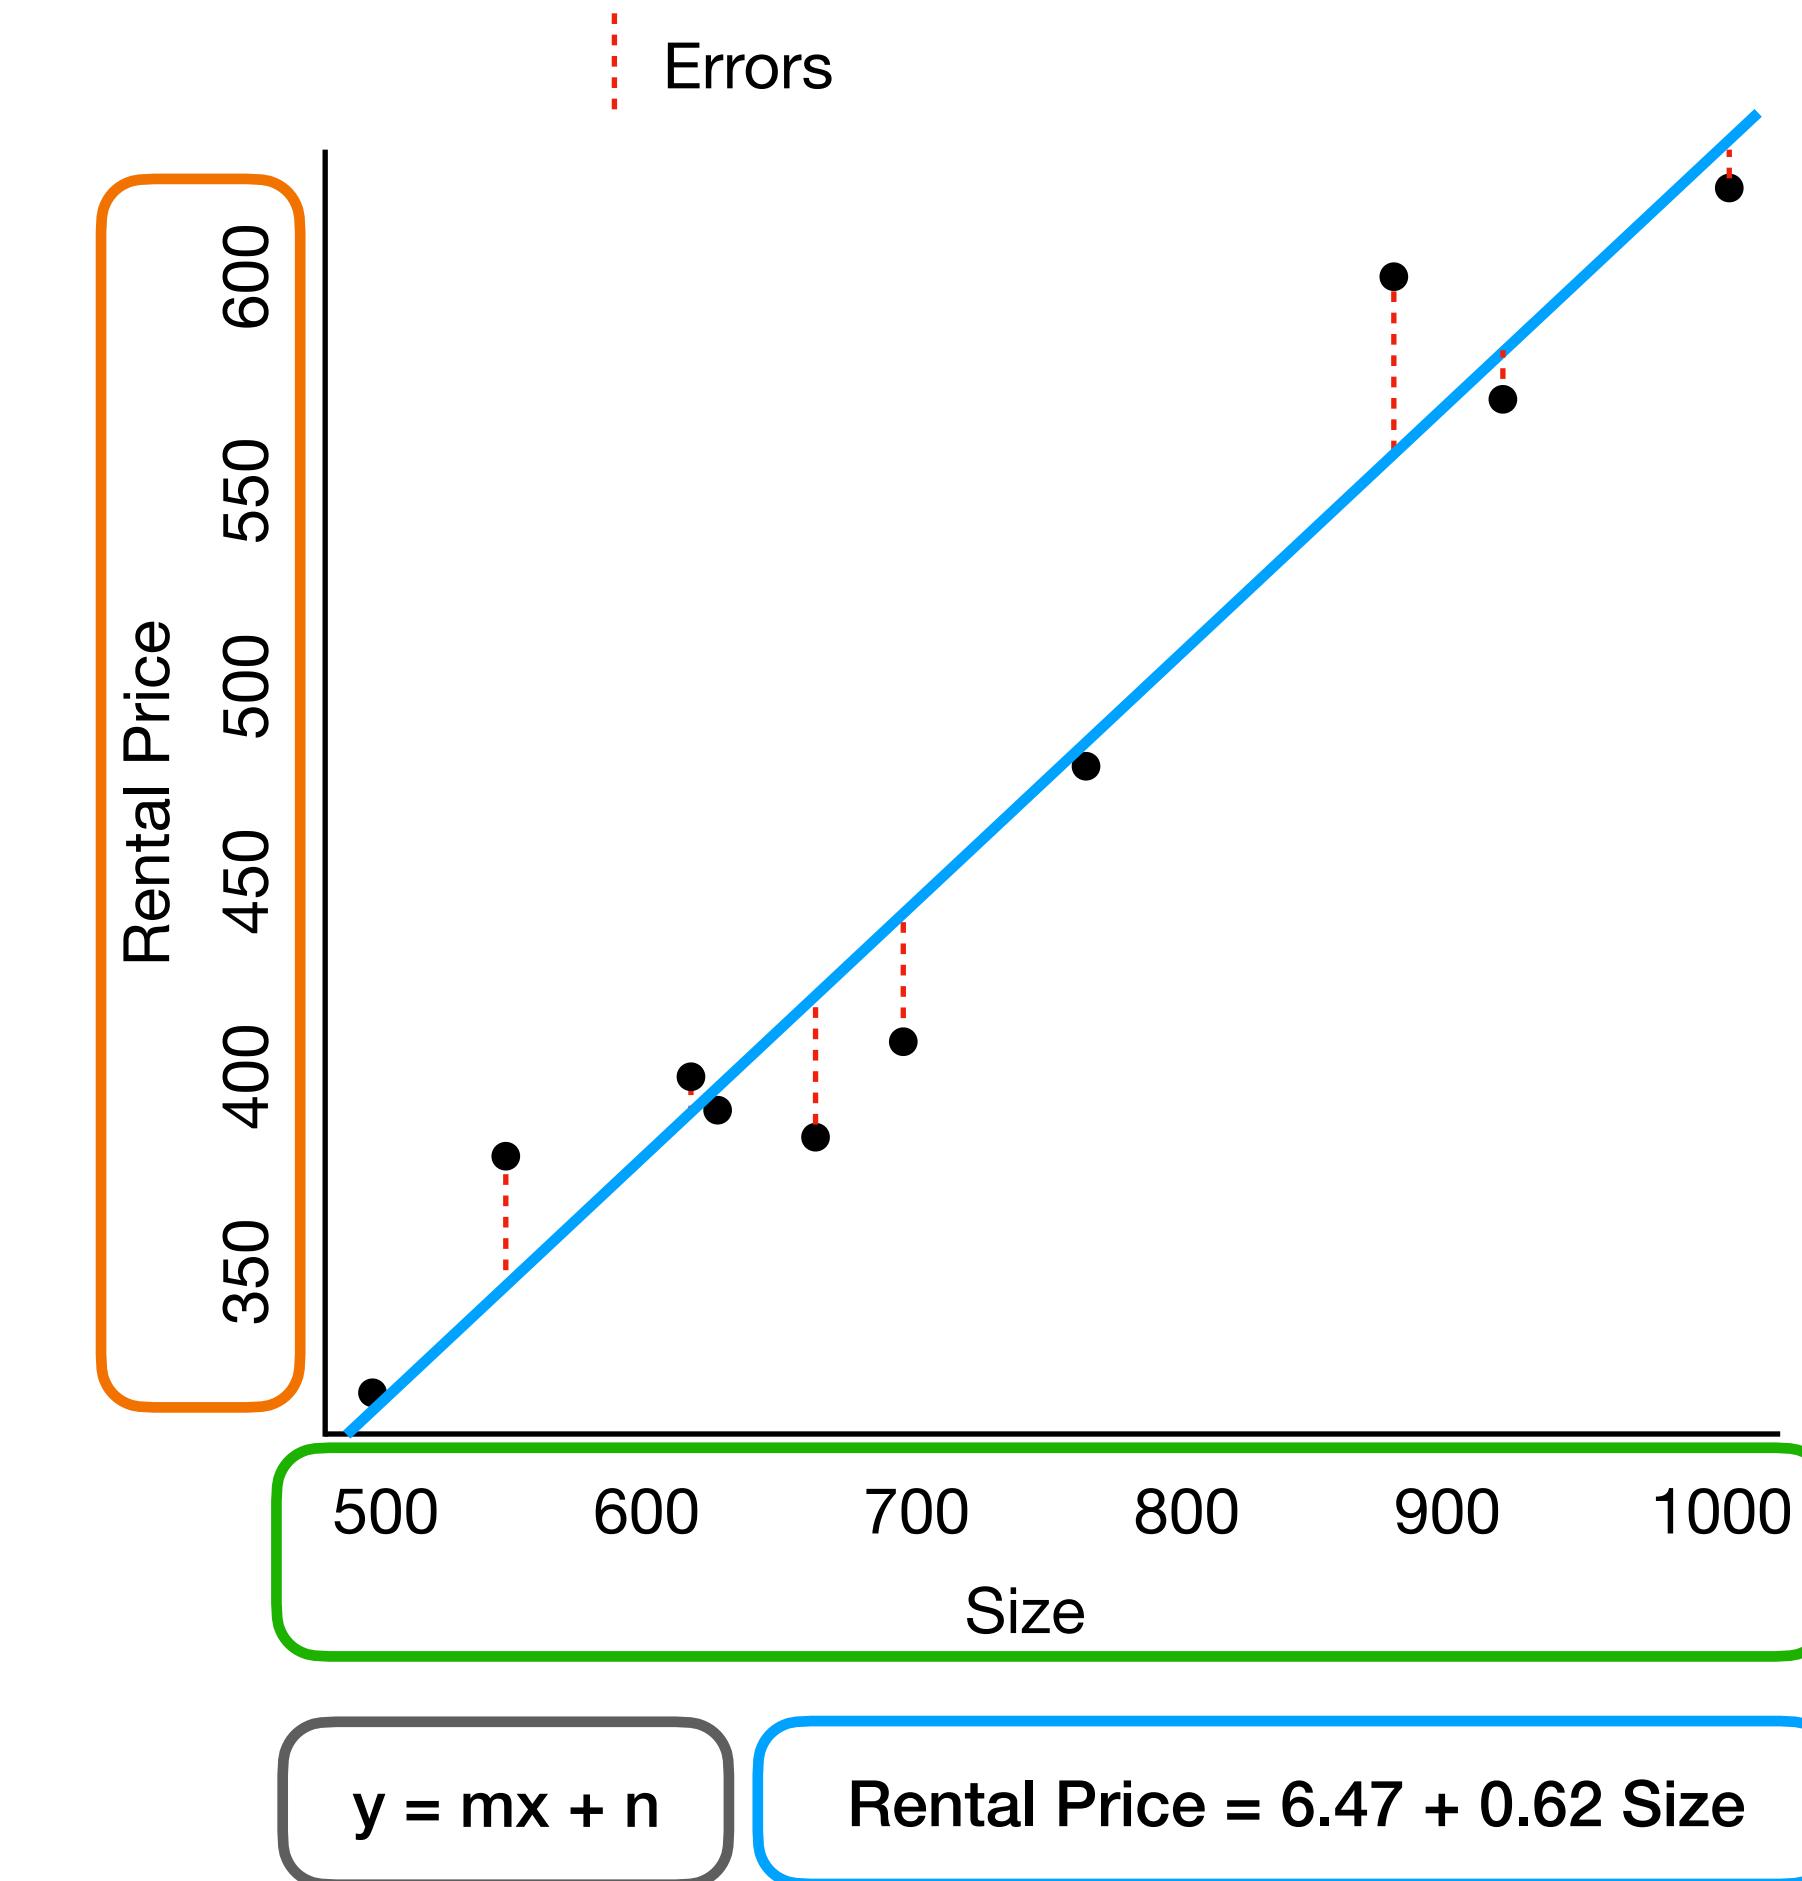
We want to obtain the line that minimizes the sum of squared errors.

$$\min(L_2)$$

To find the best parameters for the line that we are searching for (m and n in this case)

This search can be done through something called **gradient descent**

This method is at the core of many machine learning models



Linear regression model in Python to predict math scores

The technical aspects explained and many more are already handled by sci-kit learn when building any model



Load
packages

```
from sklearn import datasets, linear_model
from sklearn.metrics import mean_absolute_error
```

Therefore, our work consists of having the data ready and selecting the model. After that we are done in some lines of code

Linear regression model in Python to predict math scores

The technical aspects explained and many more are already handled by sci-kit learn when building any model



Therefore, our work consists of having the data ready and selecting the model. After that we are done in some lines of code

Load packages

Start linear model

```
from sklearn import datasets, linear_model  
from sklearn.metrics import mean_absolute_error  
  
# Create linear regression object  
regr = linear_model.LinearRegression()
```

Linear regression model in Python to predict math scores

The technical aspects explained and many more are already handled by sci-kit learn when building any model



Therefore, our work consists of having the data ready and selecting the model. After that we are done in some lines of code

Load packages

Start linear model

Fit Model to data

```
from sklearn import datasets, linear_model  
from sklearn.metrics import mean_absolute_error
```

```
# Create linear regression object  
regr = linear_model.LinearRegression()  
  
# Train the model using the training sets  
regr.fit(trainX, trainY)
```

Linear regression model in Python to predict math scores

The technical aspects explained and many more are already handled by sci-kit learn when building any model



Load packages

Start linear model

Fit Model to data

Test your model

```
from sklearn import datasets, linear_model  
from sklearn.metrics import mean_absolute_error
```

```
# Create linear regression object  
regr = linear_model.LinearRegression()
```

```
# Train the model using the training sets  
regr.fit(trainX, trainY)
```

```
# Make predictions using the testing set  
y_pred = regr.predict(testX)
```

Therefore, our work consists of having the data ready and selecting the model. After that we are done in some lines of code

Linear regression model in Python to predict math scores

The technical aspects explained and many more are already handled by sci-kit learn when building any model



Therefore, our work consists of having the data ready and selecting the model. After that we are done in some lines of code

Load packages

Start linear model

Fit Model to data

Test your model

Performance measurement

```
from sklearn import datasets, linear_model  
from sklearn.metrics import mean_absolute_error
```

```
# Create linear regression object  
regr = linear_model.LinearRegression()  
  
# Train the model using the training sets  
regr.fit(trainX, trainY)
```

```
# Make predictions using the testing set  
y_pred = regr.predict(testX)
```

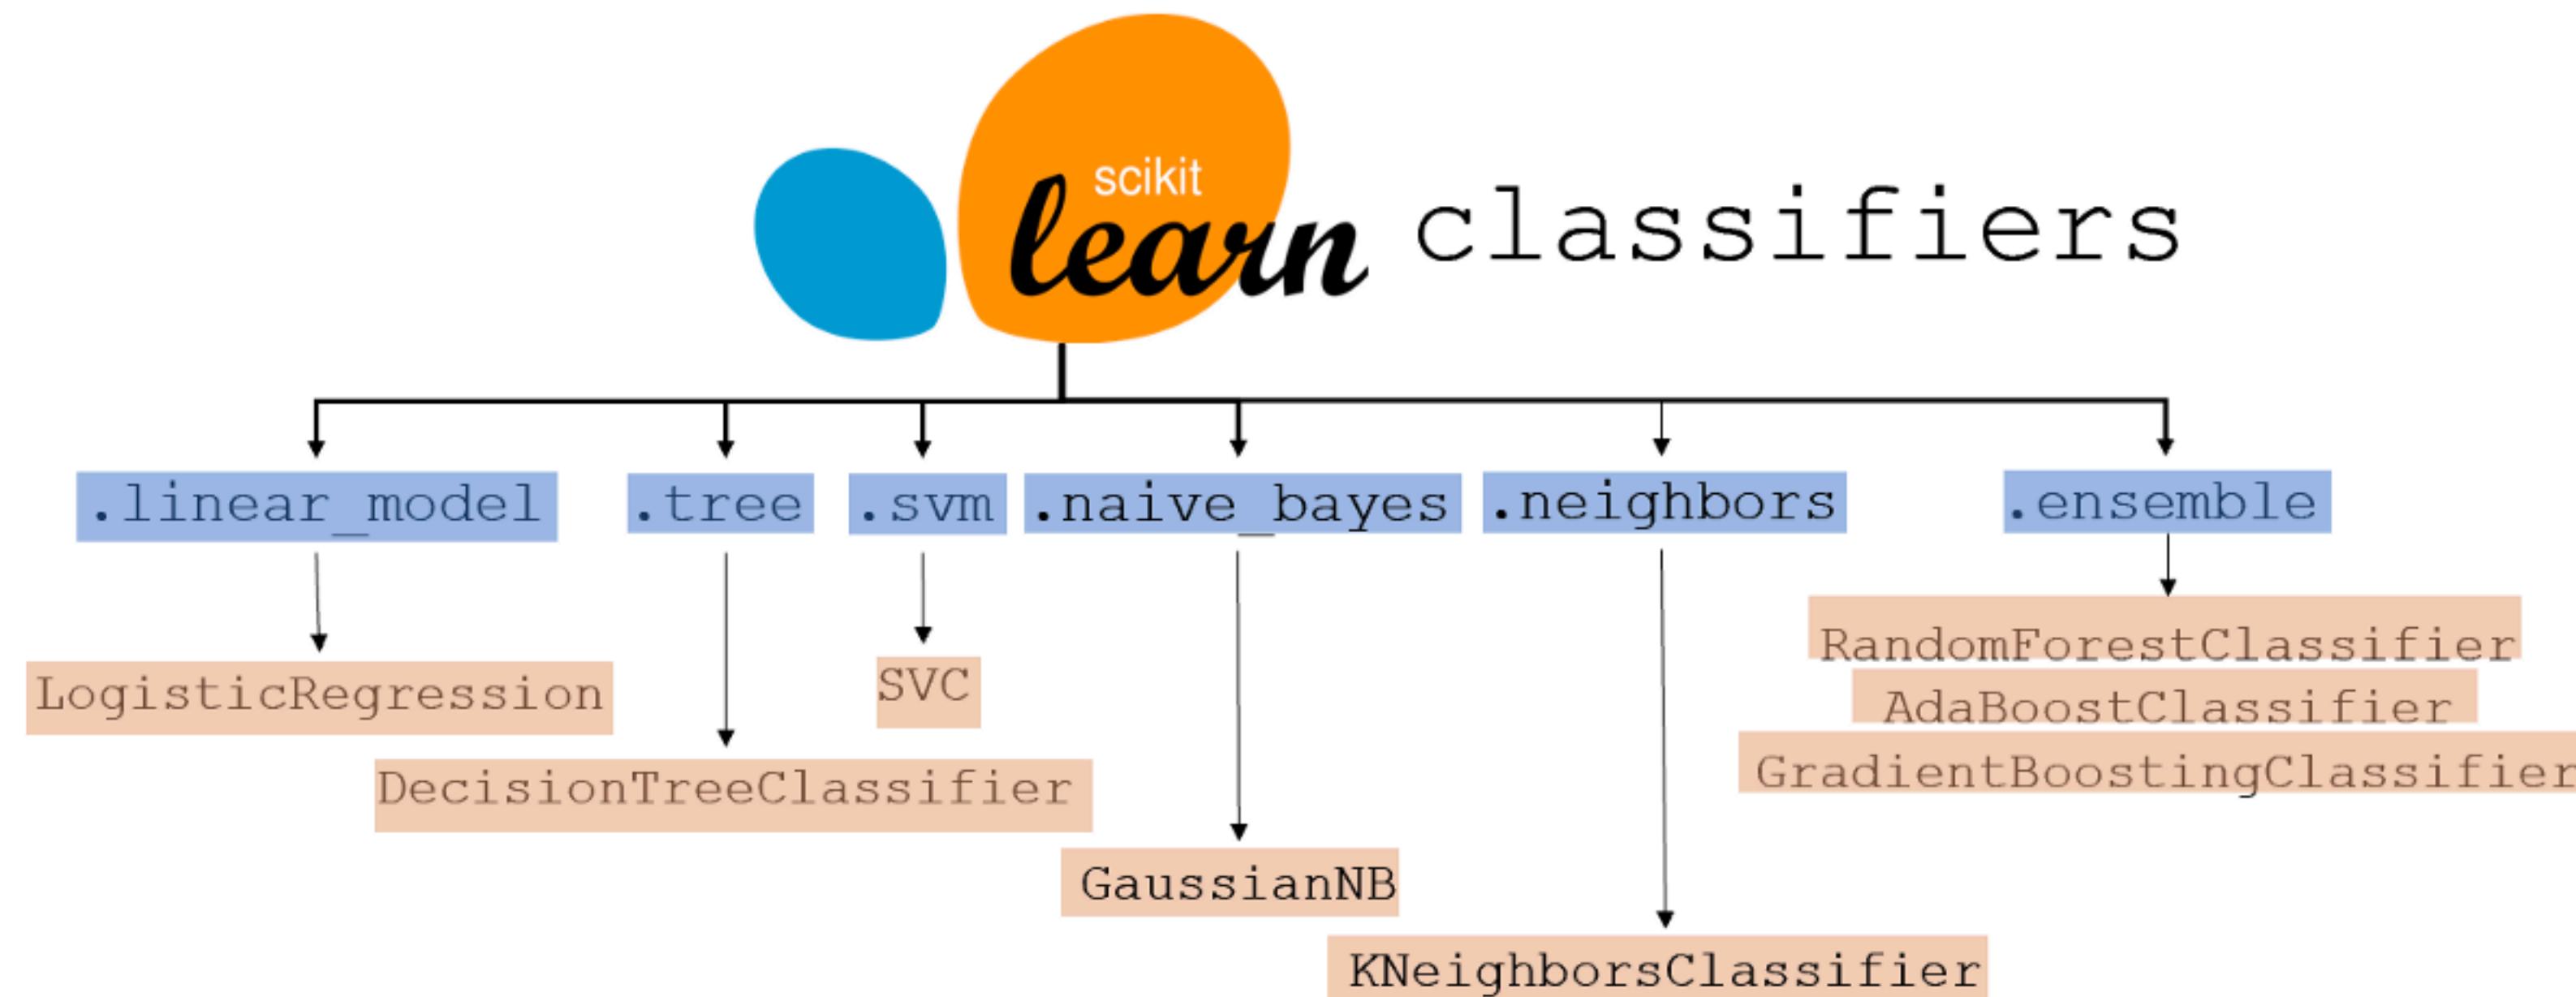
```
# The mean absolute error  
print("MAE: %.2f" % mean_absolute_error(testY, y_pred))  
MAE: 4.41
```

Exercise

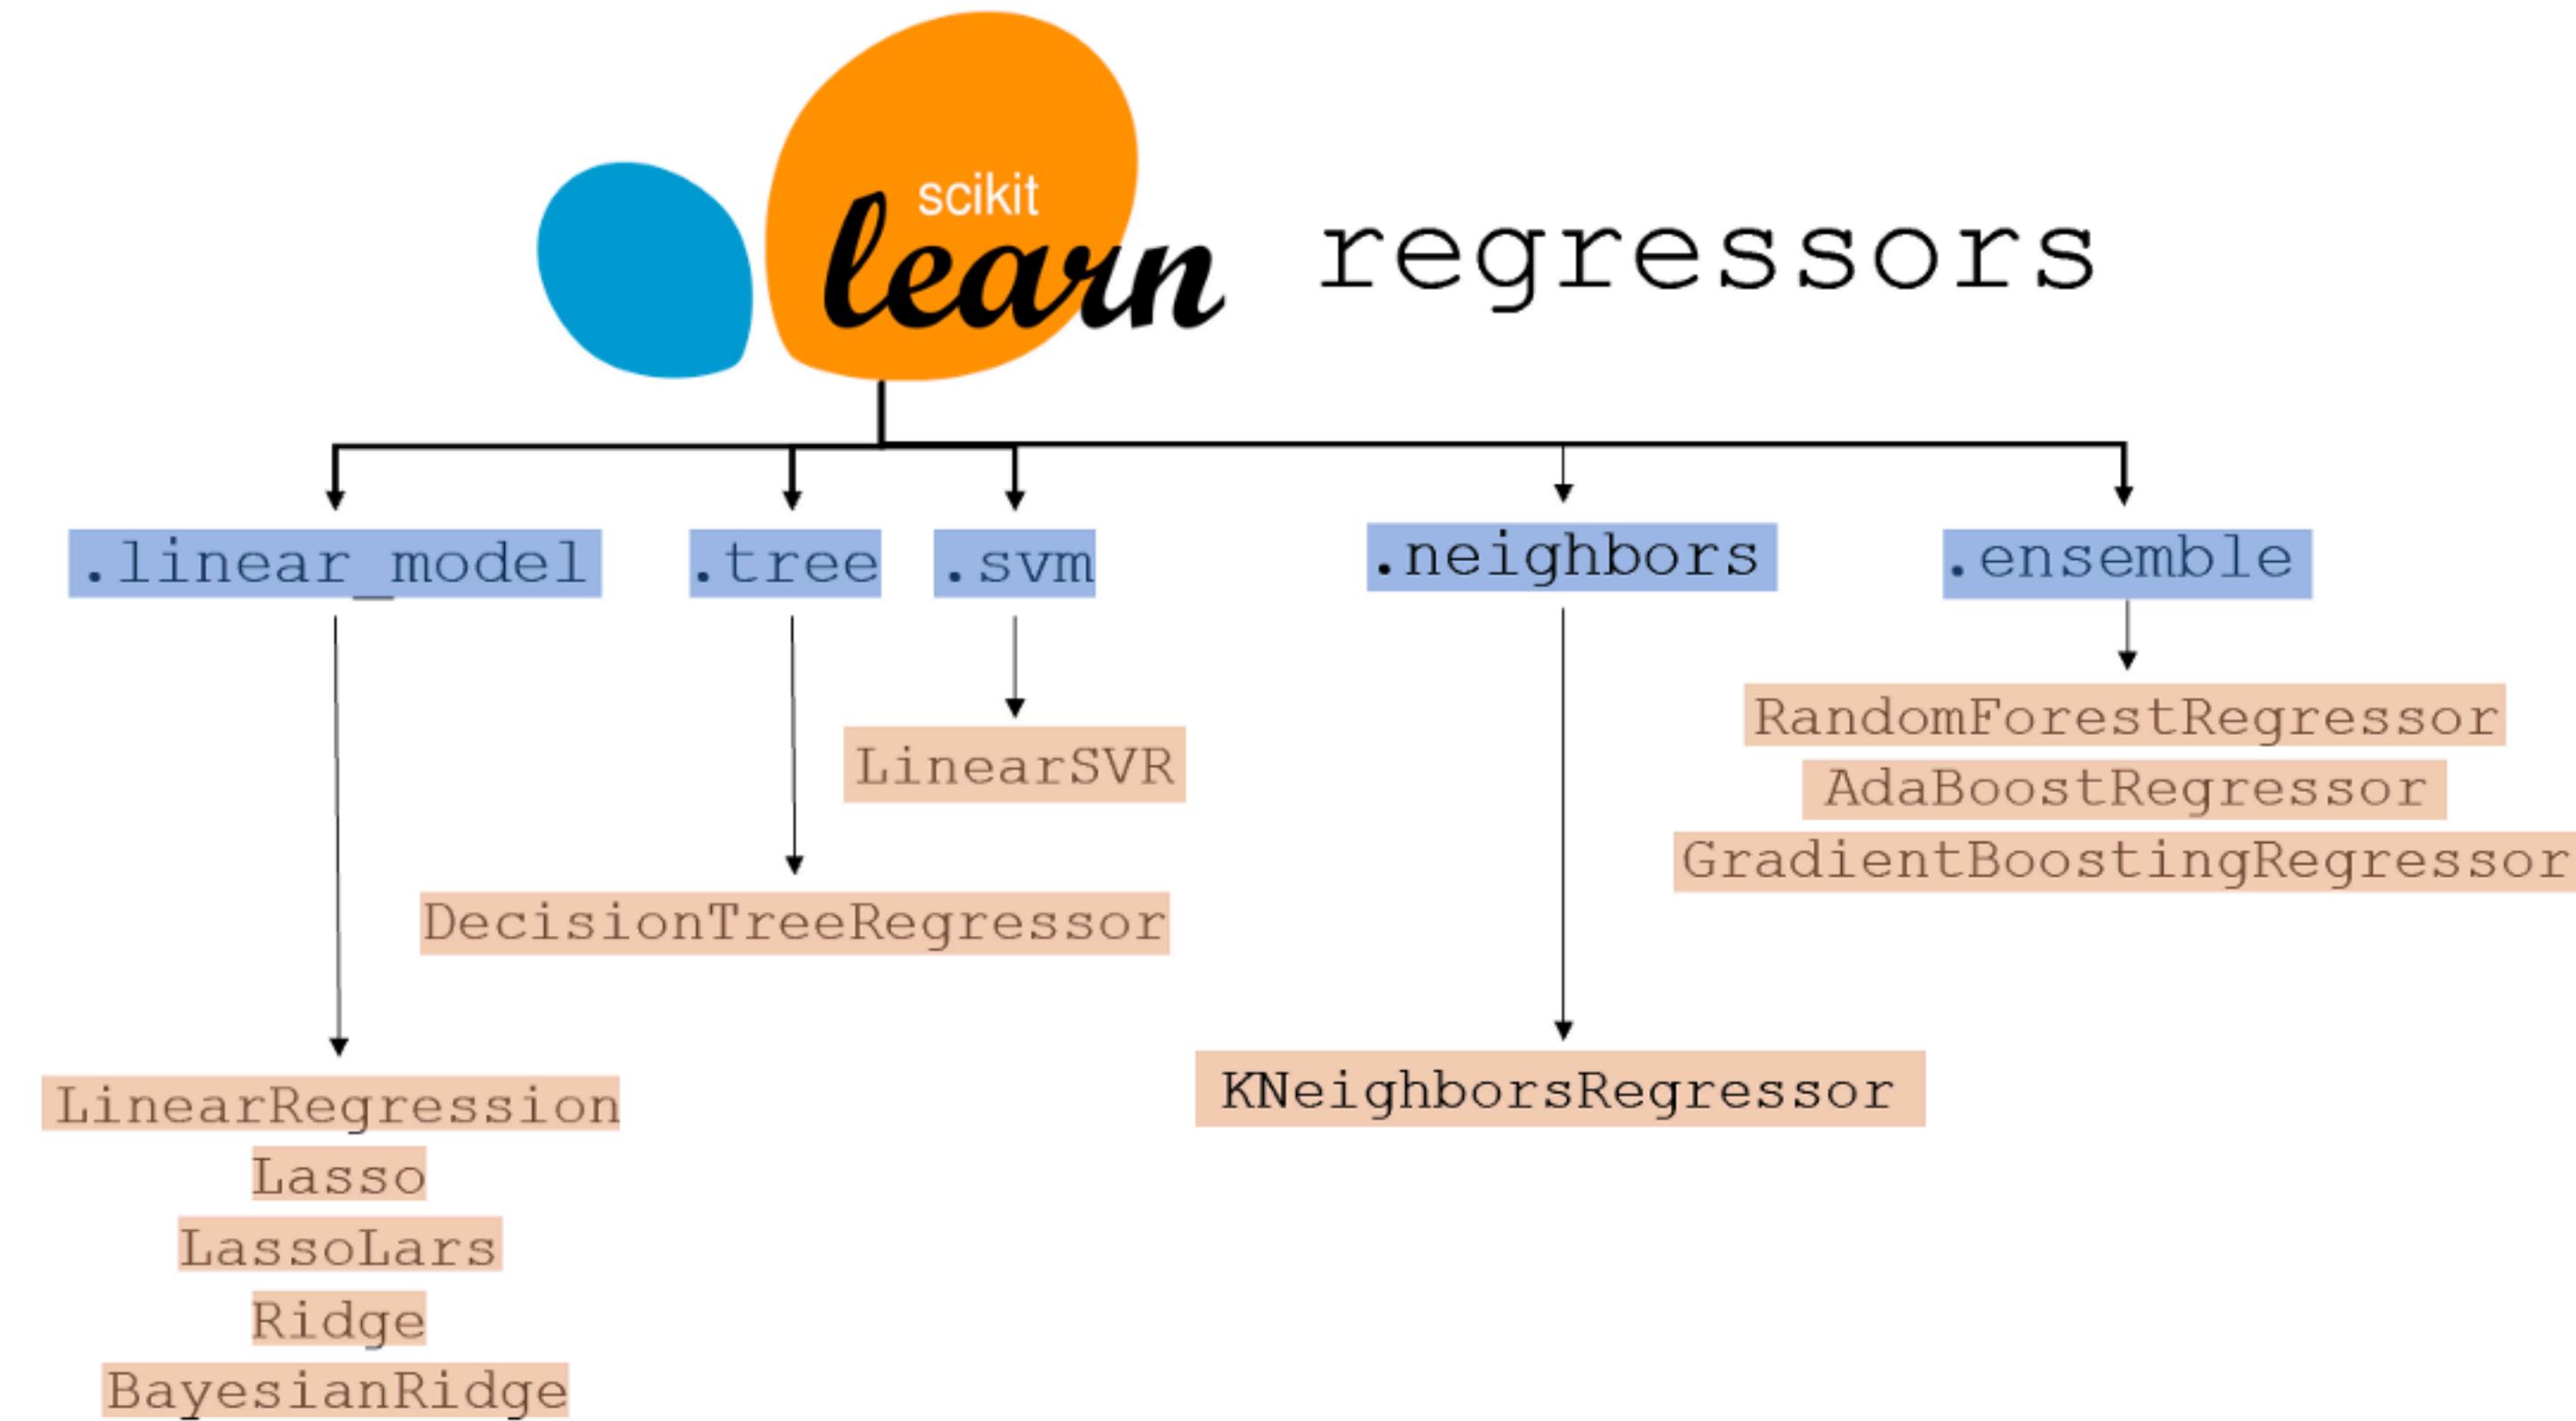
- Code your first Linear regression model and get how good it does

Topic 3: Scikit Learn models Overview and More

Scikit learn models for Classification problems



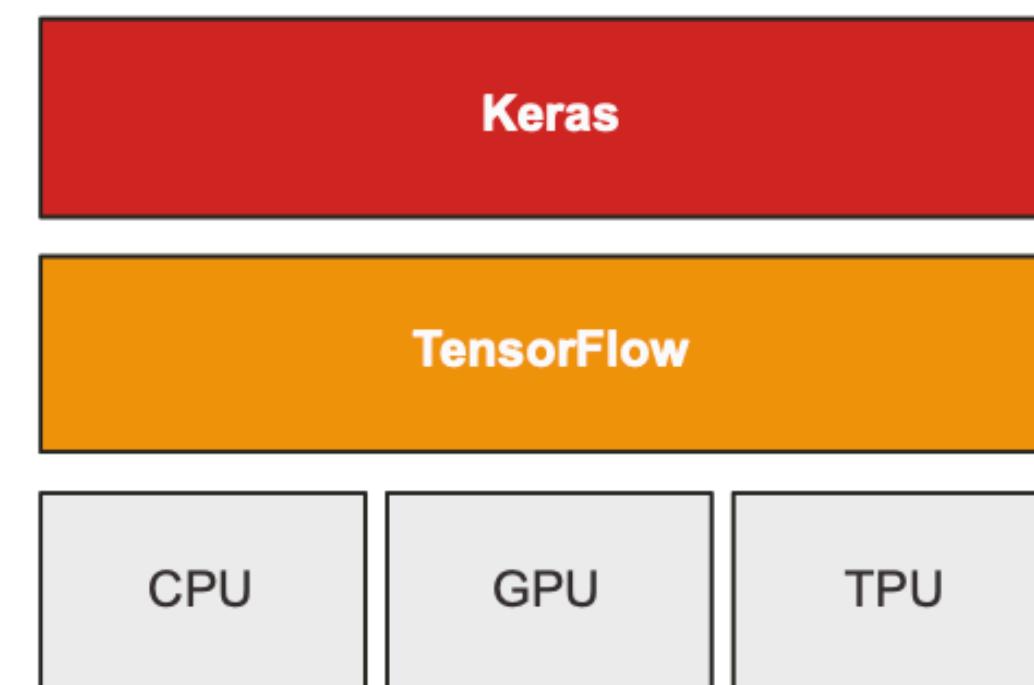
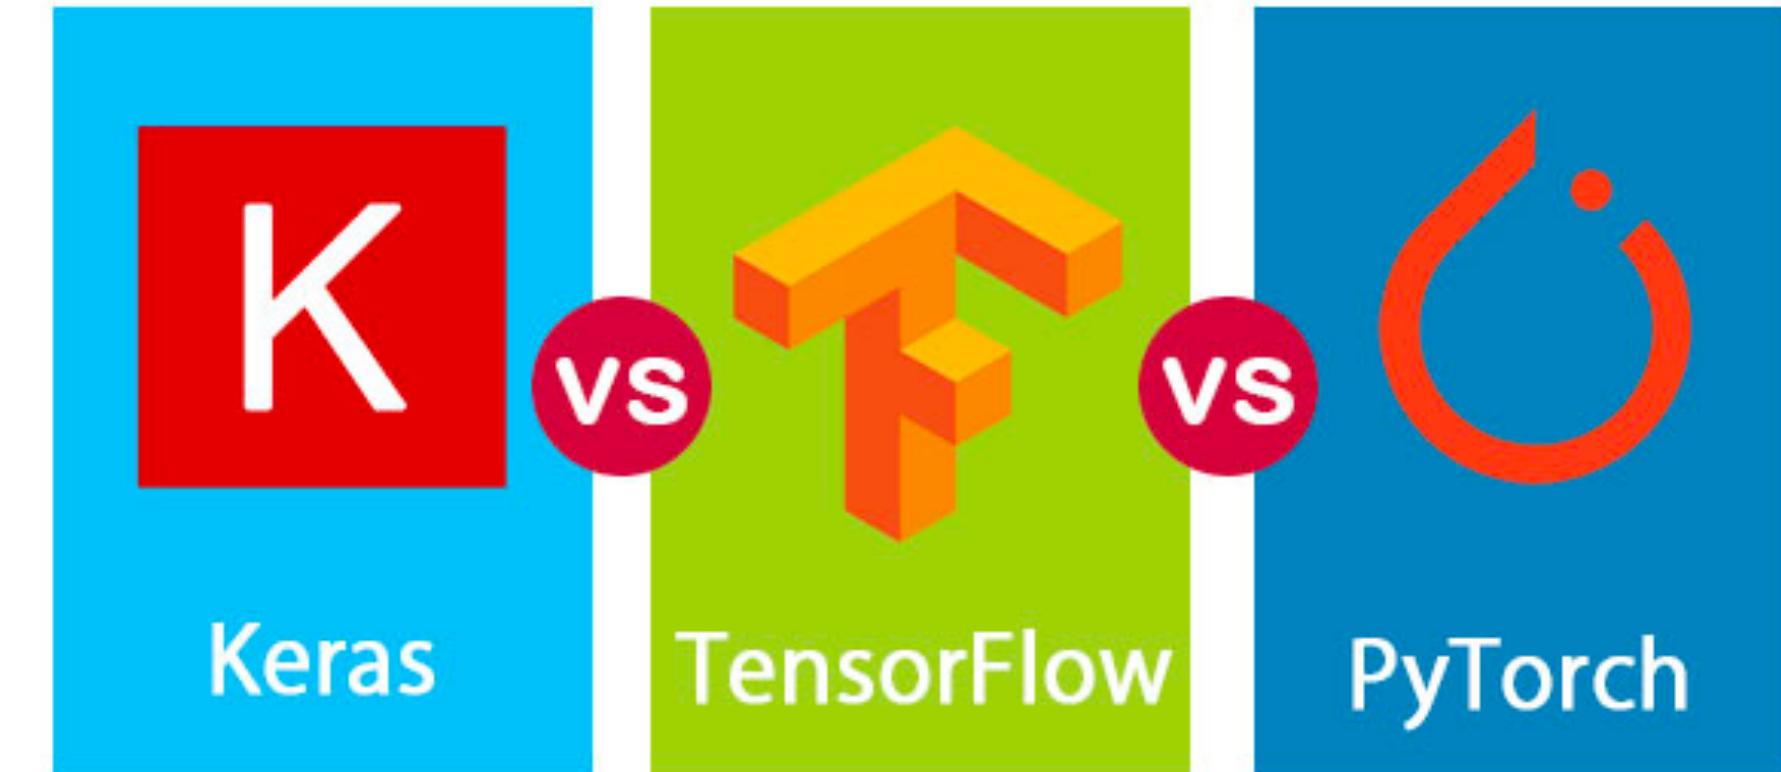
Scikit learn models for Regression problems



You can also “easily” build Neural Networks in Python



```
>>> from sklearn.datasets import make_regression
>>> from sklearn.ensemble import GradientBoostingRegressor
>>> from sklearn.model_selection import train_test_split
>>> X, y = make_regression(random_state=0)
>>> X_train, X_test, y_train, y_test = train_test_split(
...     X, y, random_state=0)
>>> reg = GradientBoostingRegressor(random_state=0)
>>> reg.fit(X_train, y_train)
GradientBoostingRegressor(random_state=0)
>>> reg.predict(X_test[1:2])
array([-61...])
>>> reg.score(X_test, y_test)
0.4...
```

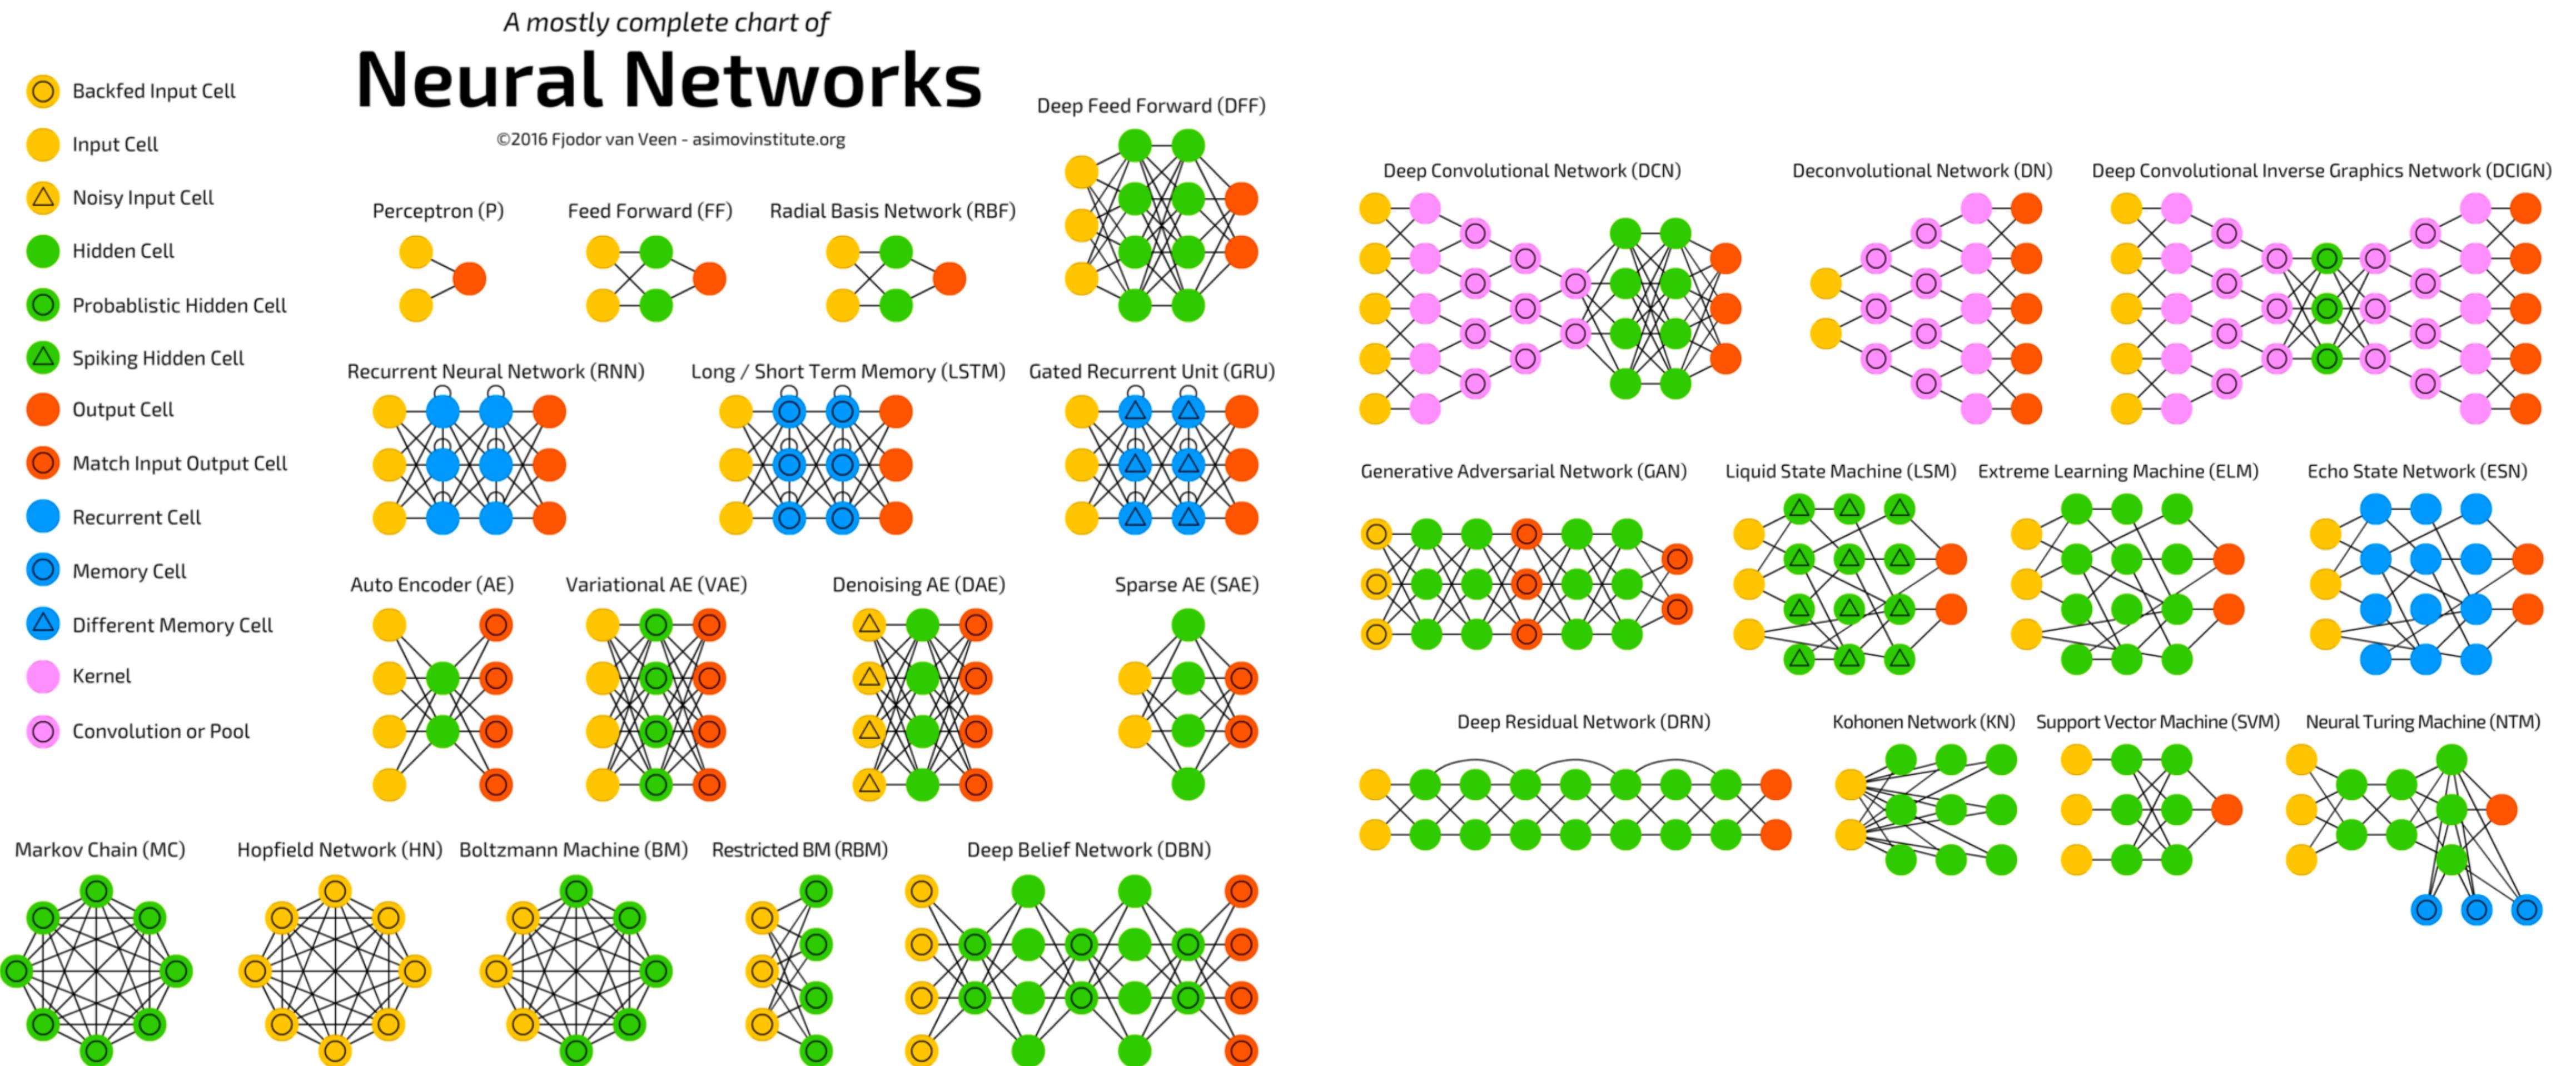


Deep learning development:
layers, models, optimizers, losses,
metrics...

Tensor manipulation infrastructure:
tensors, variables, automatic
differentiation, distribution...

Hardware: execution

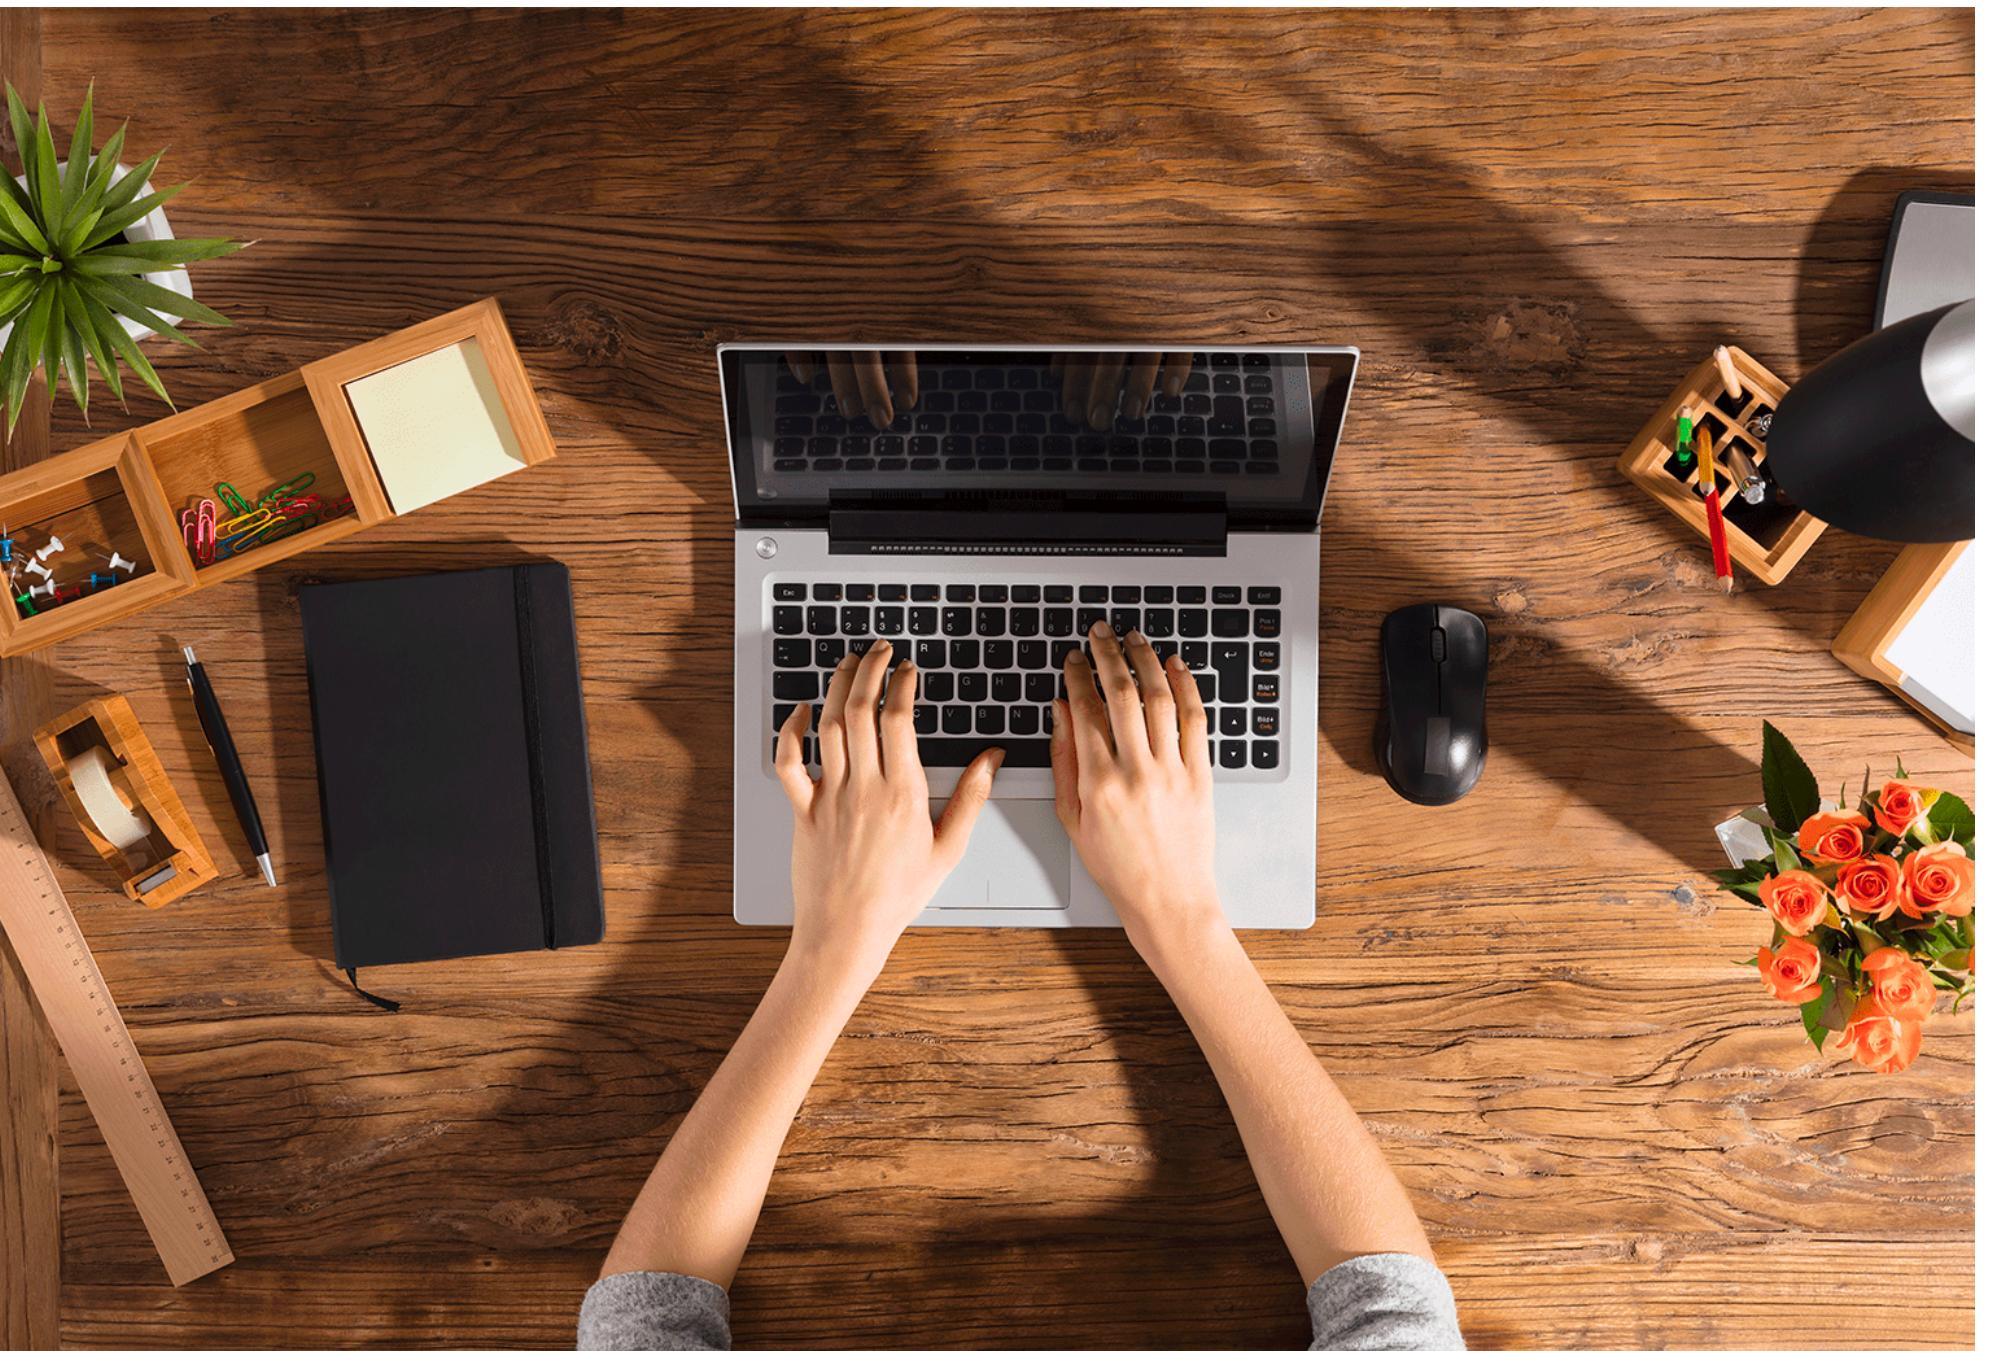
And for that you also have many options



Topic 4: Final Project Example

How is it going to be graded?

- **60% Active participation and class attendance:**
 - 15 % Attendance
 - 45 % coding exercises
 - Exercises will be uploaded to the Moodle after every session.
- **40% Final project:**
 - Python code on Jupyter Notebook, working and applying the concepts learned in class on a dataset, explaining the steps taken.
 - A 2-page report summarizing the results, the value extracted from data, and possible applications of AI.
 - Upload both to the Moodle max two weeks after the last lecture.



Final work

Code in Jupyter Notebooks

Report Explaining Results

Class 1. Intermediate Python & AI

Table of contents

- Variables
- Lists
- If-else
- For loops
- While Loops
- Functions

Variables

```
In [5]: a = 10  
        b = 5  
        print(a, b)
```

10 5

```
In [6]: a + b
```

```
Out[6]: 15
```



Study of Police deaths in the USA from 1791 to 2022

Pepe Bonet Giner
9th January 2023

- 1.- Introduction to the dataset
- 2.- Results obtained (in line with the code)
- 3.- Possible further AI applications

Final work Example (Classes 5 & 6)

Let's take a new dataset and start working as if it would be your final project

But first, do you have any questions or anything you would like me to clarify?

Documentation

Datacamp Courses

Python Fundamentals

Are you ready to gain the foundational skills you need to become a Python programmer? In this track, you'll learn the Python basics you need to start on your programming journey, including how to clean real-world data ready for analysis, use data visualization libraries, and even how to write your own Python functions.

Your instructor Hugo will introduce you to how companies worldwide use Python to gain a competitive edge. Through hands-on coding exercises you'll then learn how to store, manipulate, and explore data using NumPy. Then it's time to level-up as you learn how to visualize your data using Matplotlib, manipulate DataFrames and dictionaries using pandas, and write your own functions and list comprehension. Start this track to add these essential Python skills to your data science toolbox.

[Switch Track](#)

Python 15 hours 4 Courses 1 Skill Assessment

Data Scientist with Python

Gain the career-building Python skills you need to succeed as a data scientist. No prior coding experience required.

In this track, you'll learn how this versatile language allows you to import, clean, manipulate, and visualize data—all integral skills for any aspiring data professional or researcher. Through interactive exercises, you'll get hands-on with some of the most popular Python libraries, including pandas, NumPy, Matplotlib, and many more. You'll then work with real-world datasets to learn the statistical and machine learning techniques you need to train decision trees and use natural language processing (NLP). Start this track, grow your Python skills, and begin your journey to becoming a confident data scientist.

[Resume Track](#)

Python 88 hours 23 Courses 6 Projects 3 Skill Assessments

Free Books to learn Python & data analytics

